

NDNSF Proposal | Original and Current, Side by Side

LEFT: Complete Origin proposal (59 pages). RIGHT: Current main.pdf (50 pages).
Clean pages at original size. No red/blue text markings. September 14, 2026.

Start with a topic below or use the PDF bookmarks. Pages are grouped by topic, not aligned sentence by sentence. Boundary pages may repeat across topics. A blank panel means that version has no further page in this group; it does not mean that the topic is absent. The original and current manuscripts are unchanged.

Topic	Original PDF pages	Current PDF pages	Open at
01 Title	1	1	2
02 Abstract	2-3	2	3
03 Contents	4-5	3-4	5
04 Motivation, research gap and scope	6, 11-12, 15-16	5-8	7
05 Research questions and expected contributions	13-14	9-10	12
06 NDN background and naming	7-11	11-12	14
07 Related work and design alternatives	17-25	12-16	19
08 Architecture, responsibilities and API	26-30	17-19	28
09 Named messages and invocation workflow	30-32	20-21	33
10 Authorization, confidentiality and trust	32-34	19, 21-26	36
11 Invocation modes, streams and large objects	37-39	26-29	43
12 Collaboration plan, dependencies and execution state	29, 36-37	29-31	47
13 UAV application and validation	40-46	32-35	50
14 Distributed inference application and validation	47-52	35-40	57
15 Existing evidence and its interpretation	34-36	41-42	63
16 Proposed evaluation and metrics	45-46, 51-52, 54-55	43-45	66
17 Research plan, timeline and risks	53-56	46-47	72
18 Conclusion and expected outcomes	57	48	76
19 References	57-59	48-50	77

Baseline: reference-pdfs/Tianxing_Dissertation_Proposal_Origin.pdf, not the 11-page annotated proposal 2.pdf.
Current document checkpoint: 3b3f29f5. Full source hashes and page coverage are in main_comparison-report.json.
This copy supports visual comparison only; it does not classify semantic changes or certify academic claims.

NDNSF: A DATA-CENTRIC SERVICE FRAMEWORK FOR SECURE, COLLABORATIVE,
AND DYNAMIC NETWORKED APPLICATIONS

by
Tianxing Ma

A Dissertation Proposal
Submitted for the Committee’s Review

Academic Advisor:
Dr. Lan Wang

Major: Computer Science

The University of Memphis
June 2026

NDNSF: A DATA-CENTRIC SERVICE FRAMEWORK FOR SECURE, COLLABORATIVE,
AND DYNAMIC NETWORKED APPLICATIONS

by
Tianxing Ma

A Dissertation Proposal
Submitted for Committee Review

Academic Advisor:
Dr. Lan Wang

Major: Computer Science

The University of Memphis
September 2026

Abstract

Modern networked applications increasingly run as collections of services across mobile devices, edge nodes, cloud resources, storage systems, and users. Existing host-centric frameworks provide useful tools, but their endpoint and channel model leaves service discovery, provider selection, channel security, caching, storage, and application state in separate systems. The problem is the mismatch between this model and applications that care about named services, named data objects, and named results rather than fixed endpoint connections.

Named Data Networking (NDN) offers named and secured data retrieval, in-network caching, multicast forwarding, and data-centric validation. These primitives match service-oriented applications in dynamic networks, but they still need to be organized into higher-level service functions: invocation, authorization, provider selection, provider collaboration, and large-data exchange.

This dissertation proposes the Named Data Networking Service Framework (NDNSF), a data-centric service framework for secure, collaborative, and dynamic networked applications. NDNSF v0.1 is the completed foundation: it provides a generic service API, NDN-based request/ACK/selection/response messages, selective provider choice, admission metadata, name-based access control, and MiniNDN evaluation against NSC and gRPC.

The remaining dissertation work extends this foundation toward collaboration among service providers such as drones, edge nodes, storage services, and inference accelerators. The planned framework mechanisms include targeted invocation for known providers, trusted local service composition inside one container, and large-data references for artifacts that should be retrieved by name instead of carried as service payloads. These mechanisms are intended to be reusable framework features rather than application-specific shortcuts.

The proposal organizes the remaining framework problem into four layers: invocation, authorization, coordination, and data handling. This structure separates facts about NDN primitives from the higher-level service mechanisms that NDNSF must provide.

Abstract

A requester may know the required service without knowing which authorized Providers are currently available and willing to execute it. The initial NDNSF prototype supports discovery and selection of one Provider per invocation over Named Data Networking (NDN), but not multi-role collaboration. This proposal evaluates that foundation and investigates its extension to dependent service roles.

In the evaluated design, attribute-based encryption (ABE) protects discovery descriptors without enumerating Provider identities. ABE-backed challenge verification, signature validation under trust rules, and request-state checks govern service authorization. Service permissions remain independent of identity-signing credentials; runtime grants and withdrawals require validated permission updates and, for withdrawal, rekeying. Authorized but unselected candidates can read descriptors; request-scoped keys protect selected input and result exchanges.

The proposed collaboration extension uses eligible positive acknowledgment (ACK) Data to assign roles and derive local dependency checks. Providers are intended to accept intermediate Data only from assigned predecessors over declared edges, whether retrieved from producers or caches. The unmanned aerial vehicle (UAV) and distributed-inference applications provide validation settings for invocation, data delivery, runtime authorization, and collaboration.

Existing evaluations support limited invocation and execution-efficiency findings, not the effectiveness of the proposed collaboration mechanism. Planned evaluation compares complete authorization alternatives, tests role and dependency enforcement during execution and recovery, and measures communication, processing, and management costs. The research aims to establish when the invocation design and its collaboration extension are useful, and when their costs outweigh their benefits.

Keywords: Named Data Networking, service authorization, provider discovery, collaboration, distributed inference

Two validation workloads drive and evaluate the framework. The UAV workload tests NDN² under mobility, intermittent connectivity, command-sensitive services, video and telemetry data, and multi-UAV collaboration. The distributed AI inference workload tests NDN² under model retrieval, dependency execution, multi-provider cooperation, and large intermediate objects.

The central hypothesis is that combining NDN’s data-centric communication model with explicit service-level abstractions can reduce the complexity of building secure and collaborative networked applications. The dissertation contribution is the design, implementation, and evaluation of NDN²; the UAV and distributed inference systems validate the framework and feed requirements back into it.

Keywords: Named Data Networking, service-oriented applications, data-centric security, UAV networks, distributed AI inference, provider collaboration

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Table of Contents

Chapter	Page
1 Introduction	1
1.1 Motivation	1
1.2 NDN Building Blocks for NDN	2
1.3 Research Gap	6
1.4 Thesis Statement	7
1.5 Research Questions	8
1.5.1 Expected Answer Shape	8
1.6 Expected Contributions	9
1.7 Proposal Scope Boundary	10
2 Frameworks for Service-Oriented Applications	12
2.1 Purpose of This Review	12
2.2 Host-Centric RPC and Microservice Frameworks	12
2.2.1 What These Systems Solve	13
2.2.2 Open Issues for NDN	13
2.3 Robotics Middleware and Service Composition	14
2.3.1 Lessons for UAV Service Containers	14
2.4 Message Queues and Event Systems	15
2.5 NDN Computation and Service Frameworks	15
2.5.1 Function-Centric Computation	15
2.5.2 Serverless and Edge-Oriented NDN Frameworks	16
2.5.3 Remote Invocation Frameworks	16
2.5.4 Application-Specific NDN Frameworks	17
2.5.5 Security and Access-Control Mechanisms	18
2.6 NDN-Based Service Systems	18
2.7 Design Patterns and Trade-offs	19
2.8 Open Issues	19
3 NDN Framework	21
3.1 Design Goals	21
3.2 NDN v0.1 Baseline	21
3.3 Provider Collaboration as the Central Extension	24
3.4 Generic Dynamic API	24
3.5 Service Message Model	25
3.6 Service Invocation Workflow	26
3.7 Security Model	27
3.8 Completed Foundation: NDN v0.1 Evaluation	29
3.9 Limitations of NDN v0.1	31
3.10 Supporting Mechanisms for Provider Collaboration	32
3.10.1 Targeted Invocation	32

Table of Contents

Chapter	Page
1 Introduction	1
1.1 Service Discovery and Authorized Invocation	1
1.2 The Initial NDN Foundation	1
1.3 Application Needs and Framework Scope	2
1.4 Extending Invocation to Multi-Provider Collaboration	4
1.5 Thesis and Research Questions	5
1.6 Expected Contributions and Current Foundation	5
2 Related Work and Design Alternatives	7
2.1 NDN Background and Naming Responsibilities	7
2.2 RPC, Robotics Middleware, and Messaging Alternatives	8
2.2.1 Microservices and Edge Computing	9
2.3 Named Computation, Invocation, and Composition	10
2.4 Comparing Authorization Designs under the Same Requirements	10
2.4.1 DNMP as a Concrete Trust-Schema Design	10
2.4.2 Adapting the Alternative to NDN	11
2.5 Protected Discovery and Permission Management	12
3 Research Approach: Invocation and Collaboration	13
3.1 Components and Application Responsibilities	13
3.2 Participants, Trust, and Exclusions	15
3.3 Four Named Message Types	16
3.4 Authorization and Confidentiality	17
3.4.1 Rationale for ABE-Backed Service Authorization	17
3.4.2 Authorization Mechanisms, Lifecycle, and Costs	20
3.5 Invocation Modes and Execution Lifecycle	22
3.5.1 Discovery, Known Providers, and Local Composition	22
3.5.2 Admission, Progress, and Failure	23
3.6 Control Messages, Streams, and Complete Objects	24
3.7 Proposed Extension: From ACKs to a Collaboration Plan	25
3.8 Proposed Dependency Admission and Execution State	26
3.9 Conditional Correctness Argument	27
4 Application Scenarios and Validation Plan	28
4.1 Validation Strategy: Applications Exercise the Framework	28
4.2 NDN-UAV: Control, Sensing, and Collaboration	28
4.2.1 Application Structure and Service Workflows	28
4.2.2 Operator Workflow and State Interpretation	29
4.2.3 Representative Collaboration: Multi-View Collection and Analysis	31
4.3 NDN-DI: Inference Lifecycle, State, and Collaboration	31
4.3.1 Application Structure and Inference Workflows	31

3.10.2	Trusted Local Invocation	33
3.10.3	Large-Data Reference Abstraction	33
4	UAV Network Application over NDN	35
4.1	Why UAV Networks Need a Service Framework	35
4.2	UAV Services	36
4.3	Architecture	37
4.4	Operator Workflow and State Discipline	37
4.5	Live Video and Recording	38
4.6	Flight Control and Safety	39
4.7	Multi-UAV Patrol	40
4.8	Evaluation Plan	40
5	Distributed AI Inference over NDN	42
5.1	Motivation	42
5.2	Approaches to Distributed AI Inference	42
5.3	NDN-Based Inference Architecture	43
5.4	Automatic Collaboration Plan Generation	45
5.5	Dependency-Driven Execution	45
5.6	Current Prototype Status	46
5.7	Prototype and Evaluation Plan	46
6	Research Plan and Timeline	48
6.1	Completed Work	48
6.2	Planned Work	48
6.3	Evaluation Matrix	49
6.4	Timeline	49
6.5	Risks and Mitigation	49
7	Conclusion	52
	References	52

4.3.2	Representative Collaboration: Multi-Provider Inference	32
4.3.3	From a Model Graph to Runtime Assignments	33
4.3.4	Tensor Exchange and Autoregressive State	34
4.4	Coverage, Evidence, and Completion Boundaries	35
5	Preliminary Evidence and Proposed Evaluation	37
5.1	Measured Findings Are Configuration-Specific	37
5.2	Evidence Needed for Each Research Question	39
5.3	Application Experiments and Interpretation	39
6	Research Plan, Feasibility, and Risks	42
6.1	Core Research Scope and Feasibility	42
6.2	Limits That Must Remain Visible	43
7	Expected Outcomes and Proposal Summary	44
	References	44

Chapter 1
Introduction

1.1 Motivation

Modern networked applications are increasingly organized as collections of services rather than as single monolithic programs. A mobile robot may expose telemetry, perception, and flight-control services. A ground station may provide object detection, mapping, mission planning, and storage services. Edge devices may request AI inference services from nearby accelerators or cloud nodes. In each case, the application developer wants to invoke named services and retrieve named results without manually binding every operation to a fixed host address, a long-lived channel, and a single provider.

Today’s service frameworks provide valuable engineering tools, but TCP/IP communication is still centered on reaching endpoints and maintaining channels. Remote procedure call (RPC) systems, REST services, and microservice platforms typically identify service endpoints by host addresses, ports, or channel endpoints. Robotics middleware provides publish-subscribe or request-response communication, but often assumes an operational domain where identity, trust, network reachability, and service discovery are configured outside the service invocation itself. The problem is not that these tools are unusable; the problem is that mobile and data-heavy applications want to name the service and the data objects involved in the workflow, while the networking substrate mainly exposes where an endpoint is and how to keep a path to it. When applications become mobile, intermittent, multi-provider, and security-sensitive, developers must rebuild mechanisms for discovery, retry, authorization, integrity verification, data retrieval, and provider selection.

Named Data Networking (NDN) offers a different foundation. NDN names data directly, secures Data packets directly, and decouples retrieval from fixed endpoints. These properties naturally match service-oriented applications in which users care about a named service, a named model artifact, a named telemetry object, or a named video segment rather than a particular

Chapter 1
Introduction

1.1 Service Discovery and Authorized Invocation

A requester may know the required service without knowing which authorized Providers are currently reachable, capable, and willing to execute it. For example, an operator requesting a timely aerial observation needs a suitable camera Provider, not merely an address at which a camera service was once available. Establishing communication does not settle current willingness, permission, or task completion. The problem is this remaining application decision, rather than host-based communication itself.

Named Data Networking (NDN) supports name-based Data retrieval and data-centric security [1, 2]. Names identify Data, guide Interest forwarding, and participate in trust rules; applications can use a trust schema to constrain acceptable signing keys and guide Data validation [3]. Sync supports discovery of dataset changes so applications can retrieve available Data [4, 5]. NDN-SF builds service-level decisions from these primitives. The detailed networking background is developed in Chapter 2.

Even when one Provider suffices, candidate discovery, authorization, selection, and result retrieval remain application responsibilities. NFN, NFaaS, RICE, and NSC support named computation or invocation [6–9]. The Distributed Network Measurement Protocol (DNMP) provides a concrete example of Trust-Schema-based command authorization [10]. These NDN systems provide the computation, invocation, and authorization precedents against which this proposal examines NDN-SF’s design choices.

1.2 The Initial NDN-SF Foundation

The initial NDN-SF prototype selects one Provider from multiple candidates for each invocation. The User publishes Request Data, collects acknowledgment (ACK) Data, and publishes Selection Data for the chosen Provider, which executes and publishes Response Data.

Table 1: How NDN building blocks motivate NDNSF mechanisms.

NDN Building Block	NDNSF Mechanism
Data packets	Request, ACK, Selection, and Response are service message objects carried as named Data.
Interests	Interests fetch named Data; NDN-SVS uses Sync Interests to announce dataset updates.
Data signatures and name-based trust	Service messages, permissions, and large objects can be verified independently of the path or cache that returned them.
FIB/PIT/CS forwarding and caching	Service-related data can be retrieved by name and remain meaningful across mobility, caching, and intermittent connectivity.
NDN Sync / SVS	Users and providers discover newly published service events without binding a service to a fixed endpoint.
Segmented Data retrieval	Large models, videos, logs, and intermediate tensors are carried by reference rather than embedded in small service payloads.
NAC-ABE and named attributes	Service-level permissions and confidentiality policies are expressed over names and attributes rather than transport channels.

1.3 Research Gap

The research gap is therefore between NDN communication primitives and service-oriented application semantics. NDN provides the right lower-level building blocks: named Data, signatures, caching, forwarding state, segmented retrieval, and Sync-based discovery. A service framework must organize these building blocks into a consistent higher-level runtime.

This proposal structures the gap into four layers. At the *invocation layer*, a user needs

Candidates offer alternatives for the same invocation; the initial prototype does not assign them to different, interdependent roles.

Discovery and authorization remain research subjects independently of collaboration. A discovery descriptor supplies feasibility information, such as deadlines and resource requirements; application input supplies execution material. In the evaluated design, attribute-based encryption (ABE) protects descriptors, while protected-challenge verification, signatures under trust rules, and request-state checks govern authorization. Authorized but unselected candidates can read descriptors. The study also evaluates input/result protection, permission grants and withdrawals, and management costs.

1.3 Application Needs and Framework Scope

A networked application often needs more than the location of a service. A drone operator may require a camera near an area before a deadline, while an inference requester may require a particular model, sufficient accelerator memory, and an acceptable preparation delay. A service advertisement identifies a possible source of execution, but it does not establish that the source is willing and able to undertake this task now. NDNSF separates permission to provide a service, current willingness to participate, selection for an invocation, and completion of the requested work.

These distinctions also matter within an application that usually uses one device. After choosing a drone, an operator expects subsequent camera commands to affect that drone. Substituting another available Provider would change the meaning of the command. Conversely, an image-analysis request may permit any authorized compute Provider that meets its requirements. A reusable framework must support both situations without requiring the application to implement an unrelated communication protocol for each.

The discovery descriptor contains the information needed for a participation decision. For a sensing task, this may include a coarse area, a time window, and a sensor requirement. For inference, it may include model identity, expected input size, and a deadline. Full images,

to publish a service request, collect ACK metadata, and select one or more providers. At the *authorization layer*, users and providers need service-level permissions, name-based attributes, replay protection, and verifiable responses. At the *coordination layer*, multiple providers may need to advertise readiness, cooperate on one task, or compose local and remote services. At the *data layer*, services need to attach large objects such as video, model files, logs, and intermediate tensors by name rather than carrying them inside small request payloads.

NDNSF addresses this layered gap by treating service invocation, authorization, provider selection, collaboration, and large-data reference handling as reusable framework mechanisms rather than application-specific glue code.

1.4 Thesis Statement

This dissertation argues that dynamic service-oriented applications expose a recurring mismatch between TCP/IP’s connection-oriented model and service-oriented applications’ data-centric needs: applications want to operate on named services, named data, and verifiable results, while the network stack usually exposes addresses, channels, and endpoints. A data-centric service framework can close this gap by making service invocation, authorization, provider selection, collaboration, and large-object retrieval explicit framework-level abstractions over NDN names and verifiable Data packets.

NDNSF v0.1 is a completed core contribution of this dissertation and demonstrates the feasibility of this direction. The remaining dissertation work extends this foundation into a framework for real dynamic service applications and evaluates it through two high-pressure domains with different service requirements:

- **UAV network applications** represent mobile, command-sensitive, real-time service-container workloads, where intermittent connectivity, command authorization, telemetry, live video, recording, object detection, and multi-UAV missions must be coordinated across a ground station and multiple drones.

detailed flight instructions, and inference prompts can be delivered after selection. The boundary follows the decision: information indispensable to assessing feasibility cannot simply be withheld until after that assessment. An application should nevertheless disclose only the necessary precision. For example, an authorized camera Provider may need an approximate search area, while an unselected candidate need not obtain the full mission route. Descriptor protection limits disclosure to authorized service Providers; it does not prevent them from learning or sharing what they legitimately decrypt.

NDNSF-UAV and NDNSF-DI expose complementary demands. The former combines device-specific operations, changing observations, live data, and physical actions. The latter combines compute selection, model preparation, large artifacts, intermediate tensors, and incremental outputs. Both require authorization and bounded request lifetimes. Their differences test whether a common invocation and named-object interface remains useful beyond one workload. Reusability is an engineering objective to be evaluated through shared mechanisms and application-specific code, rather than assumed from implementing two demonstrations.

The scope is organized around an invocation’s lifecycle: deciding who may participate, deciding which work may start, and obtaining an interpretable outcome before the operation closes. These questions connect the application cases to prior invocation mechanisms [8, 9] and command authorization [10]. Table 1 groups the resulting requirements. It is a bounded scope for the workflows studied here, not an exhaustive inventory of every possible service framework. For example, billing and hard real-time flight control require additional application mechanisms. RQ1 examines participation and permission decisions, RQ2 examines dependencies when work spans Providers, and RQ3 examines observed outcomes and costs.

These are related requirements, not four strictly separated protocol layers. Authorization affects every remote operation, and execution state determines whether an otherwise valid object may be used. The following chapters explain the NDN foundation, the framework mechanisms, and their application-level validation. The research questions retain a narrower evidential purpose: determine which choices are justified, which additional guarantees the collaboration

1.7 Proposal Scope Boundary

For clarity, this proposal treats NDNFSF v0.1 as a completed core contribution of the dissertation, not as the final complete system. NDNFSF v0.1 establishes the basic architecture and feasibility of a data-centric service framework. The remaining dissertation work extends this foundation into an application-validated framework. The main framework extension is provider collaboration: how multiple authorized providers discover requests, advertise capabilities, are selected, exchange state or artifacts, and jointly complete a service task. Targeted invocation, trusted local invocation, and large-data references are treated as supporting mechanisms that make this collaborative service model practical under low-latency, same-process composition, and large-object workloads. These mechanisms are not UAV-only or AI-only features; they are proposed NDNFSF framework mechanisms whose usefulness will be evaluated through two demanding application classes.

The core dissertation question is therefore not whether two applications can be built, but whether NDNFSF can provide reusable data-centric service mechanisms that remain valid across two very different workloads.

The UAV application is used to validate NDNFSF as a service-container framework for mobile, command-sensitive, real-time systems. Its design exercises typed telemetry, mission, video, and safety information; Targeted flight-control commands; adaptive live video; encrypted local recording; object-detection service composition; and multi-drone patrol workflows. The DistributedInference application is used to validate NDNFSF as a framework for artifact-heavy, multi-stage compute/data collaboration. Its design exercises automatic model plan generation, reference-based model and activation artifacts, predictable prefetchable object names, and dependency-driven execution over tensor DAGs. The dissertation therefore does not present UAV and DistributedInference as isolated demos; it uses them as two complementary stress tests for the same data-centric service framework thesis.

These applications also contribute requirements back to NDNFSF. The UAV workload stresses low-latency command invocation, structured command status, service-container

Table 1: Framework requirements and their relationship to the research questions.

Requirement category	Required behavior	Research link
Discovery and invocation	Find willing candidates, select an executor, correlate completion or timeout	RQ1, RQ3
Authorization and confidentiality	Check both service permissions, restrict Content access, apply permission changes	RQ1, RQ3
Execution and collaboration	Associate work with active requests; enforce assigned roles and dependencies in the extension	RQ2, RQ3
Named-object delivery	Distinguish control messages, ongoing streams, and complete artifacts; validate retrieved Data	RQ2, RQ3

extension provides, and what their costs are.

1.4 Extending Invocation to Multi-Provider Collaboration

Some tasks require complementary services from different Providers. Independently invoking each service does not by itself specify which Provider fills each role or which intermediate objects a dependent service may accept. This proposal investigates an extension that makes those assignments and dependencies explicit while retaining the invocation foundation.

For example, nearby unmanned aerial vehicles (UAVs) could collect different views of a ground area for analysis by a compute Provider. The extension would assign sensing and processing roles and check each image against its declared dependency. Area and deadline information may be needed before ACK to assess feasibility. After Selection authorizes the sensing tasks, the UAVs produce images for the assigned compute Provider; Selection supplies execution details and the required access information, not images that have yet to be captured. Descriptor encryption limits reading to authorized candidates, not exposure of names or traffic patterns. This case motivates the extension; recognition algorithms remain an application choice.

composition, and large-data video references, while the DistributedInference workload stresses dependency-driven collaboration, artifact references, provider role assignment, and long-running selection diagnosis. Thus the application work has two roles in the dissertation: it is valuable in its own domains, and it is the primary method for validating and refining provider collaboration as the central NDNSE extension.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

- **Distributed AI inference** represents artifact-heavy, multi-stage compute/data collaboration workloads, where model shards, runtime artifacts, intermediate tensors, large-data objects, and heterogeneous inference providers must be selected, verified, and coordinated.

1.5 Research Questions

The dissertation is organized around the following research questions.

- RQ1. What problems in host-centric service frameworks can be addressed by data-centric service frameworks?
- RQ2. How should NDN SF expose service invocation, authorization, provider selection, collaboration, and large-object retrieval as generic framework abstractions without forcing each application to design a custom NDN protocol?
- RQ3. What reliability, security, and latency tradeoffs arise when NDN service requests are published to multiple providers and resolved through ACK metadata, provider selection, and collaboration, and how should supporting mechanisms such as Targeted invocation and compensation be used?
- RQ4. How can UAV network applications evaluate NDN SF under mobile, command-sensitive, real-time service-container workloads?
- RQ5. How can distributed AI inference evaluate NDN SF under artifact-heavy, multi-stage compute/data collaboration workloads?

1.5.1 Expected Answer Shape

The dissertation is expected to answer these questions in the following form. RQ1 should identify a set of recurring service-application problems that are not merely implementation gaps in existing systems, but consequences of binding services to hosts, sessions, or channels. RQ2 should show that NDN SF can expose service invocation, authorization, provider selection, and

1.5 Thesis and Research Questions

The thesis to be evaluated is that named, verifiable objects can support runtime service discovery and authorized single-Provider invocation, and that explicit role and dependency rules can extend this foundation to multi-Provider tasks. The study evaluates the discovery/authorization design and the collaboration extension separately. Benefits must justify their discovery, validation, communication, and key-management costs under the tested conditions.

- RQ1. How can NDN SF discover willing Providers while keeping discovery descriptors confidential from unauthorized readers and checking service permissions before execution, and at what cost relative to complete alternative designs?
- RQ2. How can the single-Provider invocation foundation be extended so Providers selected at runtime follow task roles and named dependencies, whether intermediate Data is retrieved from producers or caches?
- RQ3. Under what workload and network conditions do these mechanisms improve completion, latency, or execution efficiency, and when do their costs outweigh their benefits?

RQ1 evaluates the invocation foundation, including a DNMP-inspired role-authorization alternative paired with content encryption and adapted to the NDN SF workflow. RQ2 investigates the collaboration extension through UAV and distributed-inference dependencies; RQ3 measures both scopes and their applicable failure behavior separately.

1.6 Expected Contributions and Current Foundation

The intended contributions are an evaluated discovery/authorization design, a collaboration protocol enforcing runtime role assignments through named-dependency checks, and evidence of their benefits and costs. These build on established NDN, encryption, and computation mechanisms.

provider collaboration as reusable framework mechanisms, with local composition and large-data references serving as supporting abstractions rather than application-specific NDN protocols. RQ3 should characterize when one-to-many request publication, ACK metadata, provider selection, and provider collaboration improve reliability, and when supporting mechanisms such as Targeted invocation or compensation are needed to control latency and recovery costs without weakening security checks. RQ4 should show that UAV workloads exercise the same framework mechanisms under mobility, command sensitivity, real-time video, and multi-drone collaboration. RQ5 should show that distributed inference exercises the same mechanisms under artifact provisioning, tensor dependency exchange, provider role assignment, and long-running distributed execution.

1.6 Expected Contributions

The expected dissertation contributions are:

1. A review and gap analysis of service-oriented application frameworks with a focus on why dynamic, secure, collaborative applications still require a new data-centric framework.
2. The design and implementation of NDNSF as a generic dynamic C++ service framework over NDN.
3. Extensions to NDNSF centered on provider collaboration, with supporting mechanisms for targeted invocation, trusted local service composition, and large-data references.
4. A UAV service-container application over NDNSF, including flight control, telemetry, video, encrypted recording, object detection, and collaborative patrol services.
5. A distributed AI inference application over NDNSF, including model splitting, service plans, reference-based artifact exchange, and automatic collaboration-plan generation.
6. An evaluation of NDNSF and the application prototypes using MiniNDN, SITL-based UAV experiments, and selected real-machine validation.

Existing invocation experiments do not validate a multi-role plan. The extension adds role and dependency rules without replacing existing request-ID isolation. Later implementation work is not attributed to the initial prototype; collaboration end-to-end evaluation remains incomplete. Chapter 5 separates available findings from required evidence.

The proposal seeks committee agreement on these research questions, their significance, and a feasible method for answering them. Chapter 2 establishes the background and alternatives. Chapter 3 presents the approach, separating the invocation foundation from the proposed extension. Chapter 4 defines application scenarios and validation coverage. Chapter 5 distinguishes preliminary evidence from proposed experiments, and Chapter 6 specifies the remaining work, risks, and timeline. Completed implementation supports feasibility; it does not replace evaluation of the research claims.

IP connection. However, NDN by itself is a networking architecture, not a complete service framework. Application developers still need a runtime that defines how to publish service requests, authorize users and providers, select one or more providers, protect service data, and support higher-level collaboration.

This gap also emerged from the application-driven development in this dissertation. Building the UAV application exposed the need for stable service selection, command authorization, telemetry freshness, local service composition, and large video/recording object handling. Building the distributed inference workload exposed the need to coordinate artifact provisioning, tensor dependencies, provider roles, and large intermediate data. These experiences motivate treating the gap as a framework problem rather than as isolated application glue code.

1.2 NDN Building Blocks for NDN

This proposal assumes several NDN concepts, but they are used here for a specific purpose: to explain why NDN is a data-centric service framework. NDN is data-driven: a consumer sends an Interest packet for a named object, and the network returns a matching Data packet from a producer, cache, or other available source [1]. This basic exchange is not yet a complete service invocation, but it supplies the communication primitive that NDN builds on: the application names what it wants, not the host that should send it; the returned Data packet is independently named and signed; and retrieval can succeed even when the original producer is not on the same path as the consumer.

NDN names have three roles that are important for this dissertation. First, they identify application data, such as telemetry, video segments, model artifacts, service requests, ACKs, selections, and responses. Second, they are part of the security model: trust schemas and access-control attributes can be defined over name prefixes. Third, they guide routing and forwarding because Interests are forwarded by name prefix rather than by destination address. NDN uses the same naming foundation for service-level objects, so a service workflow can refer to both actions and large data by name.

Chapter 2
Related Work and Design Alternatives

2.1 NDN Background and Naming Responsibilities

NDN retrieves Data by name using Interest and Data packets [1, 2]. An application expresses an Interest describing the Data it seeks. A matching Data packet can be supplied by its producer or an intermediate cache. Forwarders use forwarding information to direct Interests and pending-Interest state to deliver returning Data. This permits network-layer reuse of named Data without requiring the consumer to identify its present storage location. Availability still depends on forwarding, connectivity, and a reachable copy; a name is not a delivery guarantee.

Names have several related uses in this proposal. A service name identifies the requested operation. A message name, together with the message’s validated fields, identifies an invocation exchange. An object name identifies data such as an image, model fragment, or output tensor. Names also participate in trust rules that constrain acceptable signing keys [3]. These uses should agree, but should not be conflated: matching a service prefix does not by itself prove service permission, and a valid producer signature does not prove that the object is the input assigned to the current execution.

For example, an illustrative image name could have the form

/org/uav7/view/task42/view2/version1

under the sensing UAV’s producer namespace. The application must define what the task, view, and version components mean and validate the signer. A processing Provider can refer to that exact image rather than rename it as its own observation. This example explains the naming principle, not a mandatory NDN wire format. Names and traffic patterns can remain visible even when Content is encrypted. Sensitive identifiers should therefore not be embedded in public names merely because the associated payload is protected.

The forwarding Content Store offers opportunistic caching. A repository provides an application-managed place to retain objects for later retrieval. Neither a cache hit nor a successful

NDN forwarders maintain three data structures that are important for NDN's design. The Forwarding Information Base (FIB) forwards Interests by name prefix. The Pending Interest Table (PIT) records unsatisfied Interests and aggregates multiple Interests for the same name, allowing one returned Data packet to satisfy multiple consumers. The Content Store (CS) caches Data packets for future retrieval. Together, these structures support stateful forwarding, in-network caching, multicast-style delivery, and adaptive retrieval under mobility or failures [2]. For service applications, this means that service-related data can be named and retrieved independently of a fixed address or connection. Service invocation semantics, however, still have to be built above these packet and forwarding primitives.

NDN is data-centric in its security model as well as in its naming and forwarding model. Each Data packet carries a producer signature, allowing consumers to verify integrity and producer authenticity regardless of whether the Data came from the original producer, an in-network cache, or a repository [1]. This property is central to NDN: service requests, ACKs, selections, responses, permissions, and large artifacts can all be treated as named and verifiable objects. However, signatures alone do not decide who is allowed to call a service, which provider is authorized to offer it, or which attributes should protect each service message. Those higher-level decisions are framework responsibilities.

digest check supplies the complete admission decision for a dependent service. The consumer still needs the producer's authority, the intended object identity, and the current invocation or collaboration association. Likewise, a Data packet's FreshnessPeriod controls cache freshness treatment; it is not a substitute for message expiry, current permission checks, or application replay state.

Sync helps participants learn about changes in a named dataset [4, 5]. Applications can then retrieve and validate the announced Data. Thus, publication in NDN means making a named object available and announcing it through the configured dissemination mechanism. It should not be read as an unsolicited push of a Data packet through arbitrary forwarders. Sync recovery can repair missing announcements or trigger retrieval while time remains. It cannot ensure completion after an application deadline or during an indefinitely disconnected interval.

These distinctions define the meaning of data-driven execution here. A validated Request, ACK, Selection, result, or declared intermediate object supplies evidence for a local state transition. Timers and local scheduling policies still exist. A timeout can close collection or terminate an operation, and a planning algorithm can decide which assignment to publish. The protocol-relevant decision is represented in named, verifiable messages, so remote participants do not have to infer it from an unrecorded local callback. This does not imply that all runtime state is network Data or that packet arrival alone authorizes execution.

2.2 RPC, Robotics Middleware, and Messaging Alternatives

RPC frameworks provide typed service interfaces, and gRPC supports both unary and streaming calls [11]. Robotics middleware such as ROS provides reusable application communication abstractions [12, 13]. Message-oriented middleware supports decoupled producers and consumers [14]. These are relevant alternatives to writing application-specific networking. NDN does not claim that discovery, asynchronous interaction, or distributed application composition is absent from existing middleware.

The comparison concerns how the same application requirements are expressed. An




Figure 1: Conceptual comparison of IP and NDN hourglass models, redrawn from the NDN architecture discussion [1].

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Interest Packet	Data Packet
Name	Name
Selectors order preference, publisher filter exclude filter	MetaInfo content type, freshness period
Nonce	Content
Guiders scope, Interest lifetime	Signature signature type, key locator signature bits

Figure 2: Simplified NDN Interest and Data packet structure, redrawn from the NDN architecture paper [1].

Data-driven retrieval and in-network caching are useful, but service invocation needs more than fetching a single object. Multiple providers may need to learn that a request has been published, users may need to collect ACK metadata before selecting providers, and long-running work may need status diagnosis. NDNSF therefore uses NDN Sync, especially State Vector Sync as implemented by NDN-SVS, to let users and providers discover newly published service

messages without binding the service to one fixed endpoint [3–5]. In NDN, Sync is not the service semantics itself; it is the dissemination substrate over which the framework publishes Request, ACK, Selection, and Response messages as named Data.

The implementation is built on the standard NDN software stack. NFD provides the local forwarder and face abstraction, while `ndn-cxx` provides the C++ application library for names, packets, faces, signing, validation, and segmented object retrieval [6]. These libraries allow NDN to use ordinary NDN names and packets rather than inventing a separate transport format. They also make it possible for applications to pass large artifacts by reference: models, video recordings, images, logs, and intermediate inference data can be segmented into named Data packets, signed, optionally encrypted, retrieved by name, and verified outside the small service-request payload.

For access control, NDN builds on NAC-ABE, which combines NDN naming with attribute-based encryption so that service-level authorization can be expressed over names rather than transport connections [7]. NDN uses this capability together with controller-issued permissions, service-level attributes, and per-invocation tokens. The result is that NDN supplies the data-centric communication substrate, while NDN supplies the service-layer runtime: who may call a service, who may provide it, how providers are selected, how messages are protected, and how large named objects are attached to service workflows.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Table 1: How NDN building blocks motivate NDNSF mechanisms.

NDN Building Block	NDNSF Mechanism
Data packets	Request, ACK, Selection, and Response are service message objects carried as named Data.
Interests	Interests fetch named Data; NDN-SVS uses Sync Interests to announce dataset updates.
Data signatures and name-based trust	Service messages, permissions, and large objects can be verified independently of the path or cache that returned them.
FIB/PIT/CS forwarding and caching	Service-related data can be retrieved by name and remain meaningful across mobility, caching, and intermittent connectivity.
NDN Sync / SVS	Users and providers discover newly published service events without binding a service to a fixed endpoint.
Segmented Data retrieval	Large models, videos, logs, and intermediate tensors are carried by reference rather than embedded in small service payloads.
NAC-ABE and named attributes	Service-level permissions and confidentiality policies are expressed over names and attributes rather than transport channels.

1.3 Research Gap

The research gap is therefore between NDN communication primitives and service-oriented application semantics. NDN provides the right lower-level building blocks: named Data, signatures, caching, forwarding state, segmented retrieval, and Sync-based discovery. A service framework must organize these building blocks into a consistent higher-level runtime.

This proposal structures the gap into four layers. At the *invocation layer*, a user needs

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Chapter 2

Frameworks for Service-Oriented Applications

2.1 Purpose of This Review

This chapter is not intended to be a broad survey of every service framework. Instead, it reviews what existing systems have already achieved and identifies the open issues that motivate NDN SF. The key question is not whether systems such as gRPC, ROS, microservices, or NDN Service Call are useful; they clearly are. The key question is why dynamic, secure, collaborative applications still require another framework despite the existence of those systems.

2.2 Host-Centric RPC and Microservice Frameworks

RPC systems such as gRPC [8], HTTP/REST systems [9], and microservice platforms have become the dominant way to build distributed software. They provide typed APIs, transport bindings, load balancing, service registration, and production tooling. These systems work very well when services can be identified by stable endpoints, when channels can be established and maintained, and when security can be expressed as channel security plus service-side authorization.

However, the abstraction becomes less natural for applications with mobility and intermittent connectivity. A user of a UAV video service may care about the latest frames from drone A, not whether drone A is currently reachable through a particular IP address. A distributed inference client may care about the service name for a model stage, not a fixed GPU server. A mission-planning application may need any available drone that satisfies a role, not exactly one preconfigured endpoint. Host-centric frameworks can support these scenarios, but the burden shifts to external discovery, retry, state tracking, and application-specific coordination.

digest check supplies the complete admission decision for a dependent service. The consumer still needs the producer’s authority, the intended object identity, and the current invocation or collaboration association. Likewise, a Data packet’s FreshnessPeriod controls cache freshness treatment; it is not a substitute for message expiry, current permission checks, or application replay state.

Sync helps participants learn about changes in a named dataset [4, 5]. Applications can then retrieve and validate the announced Data. Thus, publication in NDN SF means making a named object available and announcing it through the configured dissemination mechanism. It should not be read as an unsolicited push of a Data packet through arbitrary forwarders. Sync recovery can repair missing announcements or trigger retrieval while time remains. It cannot ensure completion after an application deadline or during an indefinitely disconnected interval.

These distinctions define the meaning of data-driven execution here. A validated Request, ACK, Selection, result, or declared intermediate object supplies evidence for a local state transition. Timers and local scheduling policies still exist. A timeout can close collection or terminate an operation, and a planning algorithm can decide which assignment to publish. The protocol-relevant decision is represented in named, verifiable messages, so remote participants do not have to infer it from an unrecorded local callback. This does not imply that all runtime state is network Data or that packet arrival alone authorizes execution.

2.2 RPC, Robotics Middleware, and Messaging Alternatives

RPC frameworks provide typed service interfaces, and gRPC supports both unary and streaming calls [11]. Robotics middleware such as ROS provides reusable application communication abstractions [12, 13]. Message-oriented middleware supports decoupled producers and consumers [14]. These are relevant alternatives to writing application-specific networking. NDN SF does not claim that discovery, asynchronous interaction, or distributed application composition is absent from existing middleware.

The comparison concerns how the same application requirements are expressed. An

2.2.1 *What These Systems Solve*

RPC and microservice systems solve several problems that NDN SF should not ignore. First, they give developers a familiar programming model: a service has an interface, a request type, a response type, error handling, and often a deadline. Second, production systems around gRPC and REST provide mature tooling for load balancing, logging, monitoring, deployment, and versioning. Third, channel security such as TLS gives a practical and widely deployed way to protect client-server communication. For stable cloud services, these features are often sufficient.

NDN SF does not attempt to replace this engineering knowledge. Instead, it asks which assumptions break when service providers are mobile, intermittently connected, replicated, or embedded in resource-constrained edge environments. In these settings, a service call is not only a channel operation. It is also a decision about which provider is currently authorized, reachable, capable, and worth selecting. Existing host-centric frameworks can implement this behavior, but the logic is typically distributed across service discovery systems, load-balancers, application retries, middleware plugins, and service-specific policy code.

2.2.2 *Open Issues for NDN SF*

The first issue is endpoint dependence. Even if modern service meshes hide IP addresses behind logical service names, the underlying communication still depends on reaching an endpoint through host-centric routing. The second issue is data lifetime. Once a response, video chunk, model shard, or telemetry log is stored outside the original channel, applications must add their own verification and encryption model. The third issue is provider multiplicity. RPC normally selects a single server before the request is sent; NDN SF instead allows multiple providers to observe a request, decide whether to acknowledge it, and expose metadata that helps the user choose. These differences motivate a data-centric service framework rather than another IP endpoint wrapper.

RPC experiment must account for service resolution, connection management, retries, and failure detection in the selected implementation. An NDN experiment must account for name discovery, synchronization, retransmission, validation, and Content Store behavior. A typed call is not inherently less secure than named-object exchange; security depends on the accompanying authentication, authorization, and encryption mechanisms. Similarly, network-layer caching is useful only where object reuse and the workload permit it. Reusing an image does not mean that an invocation with physical side effects may be repeated safely.

For NDN SF, the design opportunity is to connect service decisions and application data through names and verifiable objects. The evaluation must test the resulting benefits rather than equate a difference in networking abstraction with superiority. Related named-computation and command-authorization systems below provide more direct protocol precedents.

2.2.1 *Microservices and Edge Computing*

CFEC studies computation offloading across vehicular fog, edge, and cloud resources using mobility and resource information [15]. MIA-NDN studies microservice-aware Interest aggregation and PIT lifetime management based on input characteristics [16]. SECaaS studies the deployment of security functions at the multi-access edge, including access-control enforcement and anonymization [17]. These works widen the comparison beyond function invocation: resource placement, forwarding state, and deployable security functions are relevant parts of the surrounding system.

Their stated objectives should not be recast as missing NDN SF features. Placement feasibility does not establish invocation permission, parameter-aware aggregation does not determine the assigned predecessor in a collaboration, and deploying an access-control service is a different question from authorizing calls to that service. Conversely, NDN SF’s application-layer checks do not replace these systems’ resource-management or forwarding mechanisms. A claim of superiority would require a common workload and matching requirements; their published results are not interchangeable baselines for the experiments in this

2.3 Robotics Middleware and Service Composition

Robotics middleware, including ROS and ROS2 [10, 11], provides a more application-aware model for robot components. Nodes can publish topics, call services, and interact through standardized message types. ROS2 improves quality-of-service control and deployment flexibility compared with earlier ROS systems. These frameworks are important precedents for UAV applications because they treat sensors, control, telemetry, and perception as composable modules.

The remaining gap is that robotics middleware is not a general data-centric network security framework. Trust anchors, certificate deployment, data-at-rest protection, multi-provider service selection, and cross-domain service authorization are not the main abstractions. In addition, a UAV application deployed across a ground station, multiple drones, and intermittent wireless links still needs naming, verification, and recovery mechanisms that are robust outside a single local robotics domain.

2.3.1 Lessons for UAV Service Containers

ROS is particularly important for this dissertation because it shows that robotics applications are naturally service-oriented. A UAV system is not a single communication flow. It contains flight-control commands, telemetry streams, camera frames, mapping, mission planning, object detection, logging, and health monitoring. ROS-style composition validates the idea that UAV software should be decomposed into services.

However, NDNFSF targets a different layer. ROS and ROS2 are strong middleware systems for local and distributed robot components, while NDNFSF focuses on named, secure, network-visible service invocation over NDN. In the proposed UAV application, the drone app can still be viewed as a service container, but external interactions such as ground-station commands, telemetry discovery, video retrieval, and cross-drone collaboration use NDNFSF service names and data-centric verification. This distinction is important: NDNFSF is not a replacement for flight controllers or robotics middleware, but a network service framework that

proposal.

2.3 Named Computation, Invocation, and Composition

NFN studies named computation and expression evaluation; NFaaS studies execution placement in a named-function environment [6, 7]. RICE provides an invocation handshake with support for authentication and security information [8]. NSC specifies signed invocation Interests validated using an application-provided Trust Schema, and discusses protected input and result exchange [9]. These designs establish relevant invocation and authorization mechanisms. Experimental comparisons below refer to the tested NSC configuration, not every feature of its full design.

For the collaboration extension, the question is how a request gathers current willingness and feasibility information before assigning execution roles, and how that assignment constrains intermediate-object use. Long-term service permission does not identify the Provider assigned to a particular predecessor role. Neither a request ID nor a valid signature alone expresses that assignment. The proposed contribution is the enforced connection between runtime assignment and object use, rather than the manifest format itself. Service-composition work such as CaSCON [18] is a relevant comparator; the research will examine assignment timing, dependency validation, and recovery semantics alongside workflow-security and capability-based designs. Novelty beyond this proposed scope remains to be established by that comparison.

2.4 Comparing Authorization Designs under the Same Requirements

2.4.1 DNMP as a Concrete Trust-Schema Design

DNMP uses role-constrained signing to authorize measurement commands [10]. Its Network Key anchors the signing hierarchy for role keys, device keys, and the distributed schema. Given `target`, `probeType`, `probeArgs`, and the user’s `roleType`, the client finds a schema-permitted signing key among the user’s available keys. It rejects a missing match

can connect such components across dynamic NDN environments.

2.4 Message Queues and Event Systems

Message queues and publish-subscribe systems decouple producers and consumers in time [12]. They are useful for telemetry, logs, asynchronous workflows, and elastic back-end processing. This is relevant to NDN because many service applications include both request-response actions and event streams. For example, a drone may receive a command as a service request while publishing telemetry or video metadata asynchronously. A distributed inference pipeline may publish intermediate results and execution events while serving user requests.

However, many message queue deployments rely on brokers or logically centralized infrastructure. Security is often enforced by broker access control and channel encryption rather than by data objects that remain verifiable wherever they are stored or forwarded. For mobile or disconnected edge environments, broker availability and placement can become a limiting assumption. Moreover, queues do not by themselves define how a user should choose among multiple providers, verify that a provider is authorized to serve a specific service, or retrieve large service artifacts as independently verifiable data.

2.5 NDN Computation and Service Frameworks

NDN already provides useful primitives for service-oriented systems: hierarchical names, receiver-driven retrieval, in-network caching, adaptive forwarding, and data signatures [1]. A service framework over NDN must bridge the gap between these primitives and application-level service semantics. Prior NDN work explores several regions of this design space.

2.5.1 Function-Centric Computation

Named Function Networking (NFN) treats computation as a first-class named object by embedding function expressions in names [13]. This model is powerful because it allows

and selects the least privileged key if several match. The receiving Network Observer Daemon (NOD) dispatches a probe only after publication validation under the schema. Thus, this is command authorization, not merely verification of an arbitrary valid signature. DNMP also bounds publication lifetimes and retains recent publication state for replay protection, assuming loosely synchronized clocks.

2.4.2 Adapting the Alternative to NDN

For comparison, we propose retaining the NDN Request–ACK–Selection–Response workflow while replacing ABE-backed invocation checks with role-constrained signing rules. The alternative must authorize both User invocation and Provider participation, match messages to active requests, and permit execution only after Selection. These adaptations are not features attributed to DNMP itself. Both designs require descriptor confidentiality before selection, recipient-scoped input/result protection afterward, and checks for withdrawn authority.

Table 2: Alternatives under the same discovery and execution requirements.

Design	Execution decision	Non-public discovery	Management cost
DNMP-inspired alternative	Validate role-constrained signatures; check invocation state and selected Provider	Combine authorization with service-group encryption or ABE for discovery	Role credentials, trust rules, status or expiry checks, and encryption keys
NDN ABE-backed	Check signatures, RequestNonce, ProviderChallenge, current permissions, and request state	Encrypt a descriptor for a service-authorized receiver set without listing current candidates	Identity credentials, ABE policies/DKEYs, status distribution, and rekeying

Role granularity is policy-defined: a role can cover multiple actions, and credentials can be reused where the rules permit. DNMP is therefore a concrete authorization precedent, not evidence that every service needs a new signing certificate.

This comparison concerns complete designs, not measured DNMP-versus-NDN performance. The invocation-state and encryption choices must be implemented and evaluated under matched requirements before claiming a cost advantage.

computation and data retrieval to be composed within the naming system itself. It also suggests a deep integration between content names and computation names, which is attractive for in-network processing.

The trade-off is complexity. Function-centric naming can make names highly expressive, but it also makes execution semantics, trust, caching, and deployment more difficult to reason about. UAV control, live video, and distributed AI inference require explicit service authorization, provider selection, long-running state, and application-defined execution policies. NDNsf therefore chooses a service-invocation abstraction rather than a fully general function-expression abstraction.

2.5.2 *Serverless and Edge-Oriented NDN Frameworks*

NFaaS extends NDN toward serverless computing by treating functions as deployable units that may be executed at edge locations [14]. Other edge-oriented frameworks, such as CFEC for information-centric Internet of Vehicles, explore low-latency microservice execution and dynamic placement in edge environments [15]. These systems highlight an important direction: service execution should move closer to users, data sources, or resource-rich edge nodes.

The open issue is that dynamic placement is only one part of service-oriented application support. Applications also need a stable API, authorization for both users and providers, multi-provider selection, large-object handling, and clear deployment procedures. NDNsf focuses on these runtime service semantics. It can coexist with future placement mechanisms; for example, a placement algorithm could decide which drones, storage components, or inference providers should advertise a given NDNsf service.

2.5.3 *Remote Invocation Frameworks*

RICE provides remote method invocation in ICN using mechanisms such as handshakes and thanks to represent computation parameters and results [16]. NSC proposes Named Service

Recipient encryption protects selected exchanges; service-group encryption or ABE can protect discovery before participants are selected. Trust Schema alone does not provide confidentiality, but it can be combined with either mechanism.

The proposed alternative requires current revocation-status checks or certificate expiry. Deleting a certificate from a repository does not invalidate cached copies. We do not attribute NDNsf’s runtime revocation protocol to DNMP.

2.5 Protected Discovery and Permission Management

NDNsf uses ABE for two related purposes: protecting discovery descriptors before Provider selection and managing service permissions independently of identity-signing credentials. The User publishes encrypted descriptors for shared retrieval, while the Controller issues policy keys defining decryption access. The enhanced NAC-ABE library provides KP-ABE support [19]. NDNsf connects these permissions to invocation by checking protected challenge values, message signatures under trust rules, and current authorization/request state. Content encryption supplies confidentiality; the combined checks govern participation and execution. Section 3.4.1 explains the reuse of permission management and its additional costs.

The comparison must include group-key or broadcast-encryption alternatives where appropriate. A stable, small group can make conventional group encryption attractive; known Providers and simple permissions can favor signed invocation with recipient encryption. ABE becomes attractive when recipients are identified by service entitlement before their current identities are known, and multiple entitlements should remain independent of identity-signing credentials. Policy size, membership churn, encryption cost, and authority availability bound this argument.

Calls for NDN and offers an RPC-like abstraction for remote execution [17]. These systems are close to NDN’s goals because they directly address service invocation rather than only content retrieval.

The key limitation is that they do not provide the complete set of mechanisms needed by the application classes targeted in this dissertation. NSC is primarily one-to-one: a caller contacts an executor, and the protocol is not designed around simultaneous provider discovery, provider ACK metadata, or multi-provider selection. RICE provides stronger invocation semantics but requires changes to the forwarding plane. NDN keeps the framework at the application/runtime layer and uses NDN Data publication and Sync to support one-to-many service request dissemination without modifying NFD.

2.5.4 Application-Specific NDN Frameworks

Other NDN systems focus on particular domains. DNMP builds a secure network measurement framework using NDN synchronization and hierarchical topics [18]. MIA-NDN explores microservice-centric Interest aggregation for NDN [19]. SECaaS studies security-as-a-service in multi-access edge environments [20]. These projects demonstrate that NDN can be adapted to practical service-like workloads, especially when names and application semantics are designed together.

The limitation is generality. Application-specific frameworks often make assumptions about a particular data type, service domain, aggregation pattern, or deployment topology. NDN instead aims to provide a reusable framework layer that can support different application services, including UAV control, video, object detection, distributed inference, storage, and collaboration. The proposed UAV and DistributedInference applications are intended to test whether the framework is general enough for multiple domains rather than only one carefully tuned application.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

2.5.5 *Security and Access-Control Mechanisms*

NDN signs Data packets by design, but signatures alone do not answer all service questions. A service framework must decide who may request a service, who may provide a service, what data should be encrypted, how keys should be distributed, and how replayed messages should be rejected. Prior work on attribute-based encryption and NAC-ABE shows how fine-grained access control can be integrated with NDN data names [7, 21]. These ideas are central to NDNSF.

The open issue is integration with the service workflow. If access control is bolted onto individual applications, every application must separately handle permission retrieval, encryption attributes, message verification, and replay protection. NDNSF makes these part of the runtime: ServiceController distributes permissions, service messages are protected according to service-level attributes, and one-time user/provider tokens bind ACK, selection, targeted fast-path, and response messages to the current invocation.

2.6 **NDN-Based Service Systems**

NDN Service Call (NSC) [17] and other NDN application frameworks show that NDN can support service invocation without IP endpoint binding. NDN’s Interest/Data exchange, name-based routing, in-network caching, and data signatures address several issues in IP-based systems. These systems are therefore important evidence that NDN is a promising foundation for services.

The remaining gap is that existing NDN service systems have not provided a general-purpose runtime that combines:

- generic dynamic service APIs,
- controller-authorized user and provider permissions,
- one-to-many request publication,

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

- provider ACK metadata,
- application-defined selection strategies,
- adaptive admission control,
- large-data references,
- provider collaboration, and
- service containers that can be evaluated in realistic settings.

2.7 Design Patterns and Trade-offs

The related work suggests several recurring design patterns. First, service frameworks must decouple service names from provider locations. Second, large inputs and outputs should be represented as named data objects rather than embedded directly in small request payloads. Third, service state usually cannot rely only on NDN’s Pending Interest Table because service execution may be long-running, multi-step, or collaborative. Fourth, security must be part of the invocation workflow rather than an afterthought.

These patterns also expose trade-offs. Function-centric names are expressive but can be difficult to deploy. Serverless placement is flexible but requires orchestration. RPC-style invocation is easy to understand but can favor one-to-one execution. Broker-based systems decouple time but may introduce infrastructure dependencies. NDNSF’s design point is to combine an RPC-like developer API with NDN’s data-centric names, signed Data, Sync-based publication, and application-visible provider selection.

2.8 Open Issues

Table 2 summarizes the open issues motivating NDNSF.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Table 2: Open issues in existing service-oriented frameworks.

Issue	Why It Matters
Endpoint binding	Mobile services and intermittently connected providers should be invoked by name and capability rather than by stable host address.
One-to-many invocation	Reliability and collaboration often require multiple providers to hear a request and allow the user to select one or more of them.
Data-centric security	Data should remain signed, verifiable, and optionally encrypted at rest, not only while moving through a secure channel.
Provider authorization	A user must verify not only that it is allowed to call a service, but also that the selected provider is authorized to offer it.
Provider selection	Provider metadata such as load, queue size, location, and availability should be part of the service workflow.
Large data	Models, video, images, logs, and intermediate tensors should not be forced into small request payloads.
Collaboration	Tasks may require multiple providers, compensation requests, or aggregation of partial results.
Operational validation	Real systems need repeatable certificate, routing, configuration, and experiment setup to validate the framework outside toy examples.

NDNSF is designed to address these issues using NDN names, verifiable Data, sync-based publication, and service-level authorization.

This chapter supports RQ1 by turning the related-work review into a set of framework requirements rather than a survey checklist. It also prepares RQ2 and RQ3 by identifying which properties must be provided by the NDNSF runtime instead of being reimplemented separately by each UAV or inference application.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Chapter 3
NDNSF Framework

3.1 Design Goals

NDNSF is designed around five goals.

- 1. **Application-facing simplicity.** Developers should register and invoke typed services without writing service-specific NDN protocols.
- 2. **Data-centric security.** Service messages and permissions should be signed, verifiable, and encrypted where appropriate.
- 3. **Dynamic provider selection.** A request should be able to reach multiple providers and select the best one according to strategy.
- 4. **Collaboration and resilience.** The framework should support multiple providers, compensation after partial failure, and cross-validation.
- 5. **Operational validation.** The framework should support MiniNDN evaluation, selected real-machine experiments, certificate management, and reproducible configuration.

3.2 NDNSF v0.1 Baseline

NDNSF v0.1 includes:

- a generic dynamic C++ service API;
- ServiceController, ServiceUser, and ServiceProvider runtimes;
- service request publication over NDN synchronization;
- Request, ACK, Selection, and Response message flow;
- user and provider permission distribution;

Chapter 3
Research Approach: Invocation and Collaboration

This chapter first describes the invocation and authorization design under evaluation, then introduces the proposed collaboration extension. Later confidentiality and lifecycle refinements are not attributed to every initial-prototype path. The basic invocation selects one Provider; role assignments and named dependency edges belong to the extension.

3.1 Components and Application Responsibilities

The framework separates application handlers from the machinery that discovers Providers, correlates messages, and checks invocation state. A ServiceUser manages outgoing requests and their callbacks. A ServiceProvider registers service handlers and processes selected invocations. A ServiceController manages permissions and authorization material. These are software roles: a ground station can consume camera services while providing image analysis, and a compute node can request an upstream service while executing its own assigned work.




Figure 1: Responsibility groups in the framework. Lines denote software dependencies, not packet directions or a mandatory deployment topology. The Controller is a logical authority, not an on-path relay for every exchange.

Application code defines the meaning of a service and supplies its feasibility and execution logic. The framework cannot determine from a valid ACK whether a camera truly

- NAC-ABE-based service-level access control;
- selective ACK and provider metadata support;
- built-in provider-selection strategies;
- adaptive admission control;
- an NDN-SVS patch used by the evaluation path;
- Python bindings for experiment and workload integration;
- MiniNDN evaluation against NSC and gRPC.

This baseline is substantial because it already connects the main pieces that a data-centric service framework must coordinate: service naming, service-message publication, user and provider authorization, encrypted service payloads, provider metadata, provider selection, and response verification. In other words, NDN v0.1 is not only an API sketch; it is an executable runtime that shows how NDN names, Sync, signed Data, NAC-ABE attributes, and service-level policy can be combined into one service-invocation workflow. This is why the proposal can focus on extending and validating the framework rather than starting from a blank design.

A paper describing this completed foundation has been submitted to ICNP. The proposal therefore treats NDN v0.1 as prior completed dissertation work and uses the remaining chapters to define how that foundation will be extended and validated through application-driven workloads.

Figure 3 summarizes the three runtime roles implemented in the baseline. The controller manages policies and distributes permissions, providers subscribe for authorized requests and publish service results, and users publish requests, collect ACK metadata, select providers, and verify responses. This role separation is the foundation for the later UAV and distributed-inference service-container work.

has an unobstructed view or whether an advertised GPU can meet a prediction deadline. It can associate the ACK with the request, verify the participating identity and required protected values, and apply the configured selection policy. A positive ACK expresses willingness subject to the request’s conditions; it is not evidence that execution has finished or a reservation unless the protocol explicitly provides one.

The ServiceContainer supports composition of runtime roles and local services within a process. Its local registry gives trusted components a way to invoke registered application functionality without turning every internal call into a network exchange. The application must keep the runtime and callback dependencies alive until active operations are closed. Local composition is useful for shared camera or inference facilities, but remote callers cannot select it as an alternative authorization path.

The current typed C++ interface illustrates this separation. The following shortened example uses application-defined, serializable status messages; identity provisioning, event-loop setup, and error handling are omitted. It illustrates API use rather than a standalone executable.

Listing 3.1: Typed registration and invocation of a compact status service.

```
provider.addHandler<StatusQuery, StatusReply>(
    ndn::Name("/status"),
    [] (const ndn::Name& requester,
        const StatusQuery& query, StatusReply& reply) {
        readApplicationStatus(query, reply);
    });

user.RequestService<StatusQuery, StatusReply>(
    candidates, ndn::Name("/status"), query,
    [] (const StatusReply& reply) { displayStatus(reply); },
    [] () { reportTimeout(); }, 2000);
```

The handler interprets the query and fills the reply. The User interface supplies an




Figure 3: NDNF v0.1 runtime roles used as the proposal baseline.

invocation timeout and result callbacks, while NDNF manages the surrounding exchange. A service requiring confidential execution input uses the protected input-delivery path described below; this compact example does not prescribe placing arbitrary application input in discovery. Typed serialization removes some repeated application glue, but does not establish semantic compatibility between different service versions. That remains part of the service contract.

3.2 Participants, Trust, and Exclusions

These are protocol roles, not necessarily separate devices: one node may be both a User and a Provider. The managed organization registers identities and approves service permissions; the Controller administers that policy. Participants require a provisioned trust anchor and validation rules for identity certificates and Controller material. Not knowing which candidates will respond to this invocation does not mean that the authorization authority has never enrolled those entities. Billing, arbitrary business transactions, globally optimal scheduling, and protection against all denial-of-service traffic are outside the claimed scope. The User requests a service and selects its Provider; in the collaboration extension, it also coordinates role assignments. Providers offer and execute services. A Controller administers authorization material and signed status. Forwarders and caches can deliver Data without being trusted to authorize execution. The model trusts the Controller and the User's planning decision; it does not claim consensus against a malicious coordinator. An attacker may observe visible names and traffic, inject invalid objects, replay previously valid messages, or operate without the needed service permission. Credential compromise and intentional sharing by an authorized principal require additional controls and are not prevented by ABE alone.

Authorization is permission to perform an operation, authentication establishes the signing identity under trust rules, and confidentiality limits who can recover Content. Valid credentials do not prove honest load reports, correct inference, or truthful sensor observations. Signed status must be obtained and validated before it can affect local authorization; a disconnected participant cannot be assumed to learn revocation immediately. Deadline/expiry

3.3 Provider Collaboration as the Central Extension

The central extension beyond NDNSTF v0.1 is provider collaboration. In a collaborative service, one logical request may involve multiple authorized providers rather than a single endpoint. Providers may advertise ACK metadata such as availability, queue state, location, or role capability; the user may select one or more providers after an ACK collection interval; and the application may issue compensation requests when only part of a logical task is completed. Long-running selected work should also be diagnosable: after a timeout, the user should be able to learn whether a selected provider has not seen the selection, is queued, is executing, has failed, or has already published a response.

This collaboration model is intentionally application-independent. Multi-UAV patrol can use it to assign mission segments and recover from partial failure, while distributed inference can use it to assign model roles and coordinate intermediate artifacts. In both cases, NDNSTF should provide the generic service semantics: request publication, provider metadata, provider selection, authorization, status diagnosis, and response verification. The application defines the domain-specific work; the framework defines how selected providers participate in one logical service workflow.

3.4 Generic Dynamic API

NDNSTF v0.1 moved toward a generic dynamic API instead of generated, service-specific stubs. A provider can register a typed service handler:

```
provider.addHandler<RequestT, ResponseT>(
    ndn::Name("/ObjectDetection/YOLOv8"),
    handler);
```

A user can invoke a service through the generic runtime:

```
user.RequestService<RequestT, ResponseT>(
    providers,
```

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

```
ndn::Name("/ObjectDetection/YOLOv8"),
request,
onResponse,
onTimeout,
timeoutMs,
strategy);
```

This API is intended to keep application data in application-defined request and response objects while the framework handles naming, publication, permission checks, tokens, provider selection, and response delivery.

3.5 Service Message Model

NDNSF v0.1 uses four framework messages to separate service publication, provider admission, provider selection, and result delivery. This message model is central to the one-to-many service design: a request can be observed by multiple authorized providers, providers can expose availability and metadata before execution, and the user can explicitly bind the request to selected providers.

Table 3: NDNSF v0.1 service messages.

Message	Producer	Purpose
RequestMessage	User	Publishes a named service request and application payload.
RequestAckMessage	Provider	Reports whether the provider is willing to serve the request and exposes runtime metadata for selection.
SelectionMessage	User	Selects one or more providers and binds the request to the selected provider tokens.
ResponseMessage	Provider	Returns the service result from a selected provider.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

```
ndn::Name("/ObjectDetection/YOLOv8"),
request,
onResponse,
onTimeout,
timeoutMs,
strategy);
```

This API is intended to keep application data in application-defined request and response objects while the framework handles naming, publication, permission checks, tokens, provider selection, and response delivery.

3.5 Service Message Model


NDNSF v0.1 uses four framework messages to separate service publication, provider admission, provider selection, and result delivery. This message model is central to the one-to-many service design: a request can be observed by multiple authorized providers, providers can expose availability and metadata before execution, and the user can explicitly bind the request to selected providers.

Table 3: NDNSF v0.1 service messages.

Message	Producer	Purpose
RequestMessage	User	Publishes a named service request and application payload.
RequestAckMessage	Provider	Reports whether the provider is willing to serve the request and exposes runtime metadata for selection.
SelectionMessage	User	Selects one or more providers and binds the request to the selected provider tokens.
ResponseMessage	Provider	Returns the service result from a selected provider.

checks require an explicit clock policy. Tests must report the allowed skew and failure behavior rather than silently assume synchronized clocks.

3.3 Four Named Message Types



Controller: authorization material and signed status.
Sync: discover names. Interests: retrieve Data.
User/Provider: validate Data, check permissions and request state before acceptance.

Figure 2: Logical invocation dependencies. Arrows express application causality, not unsolicited network Data delivery.

Table 3: Logical message responsibilities; this is not a wire-format listing.

Data	Purpose	Principal checks
Request	Publish service requirements, minimized discovery descriptor, request identity, and RequestNonce	Validate signature and current authorization context; recover the protected descriptor before evaluating willingness
ACK	Report positive or negative participation decision; a positive ACK carries selection information and ProviderChallenge	User checks signer under trust rules, service permissions, request ID and returned RequestNonce; only eligible positive ACKs enter Provider selection
Selection	Notify the chosen Provider and supply protected input or input references and request-scoped keys	Provider validates User signature and Selection using stored ProviderChallenge; checks intended Provider, request state, and service permissions
Response	Publish completion/result metadata or references to protected results	Check authorized producer, request context, result identity, and decryption scope

Table 4: Representative fields carried by the four message types.

Message	Representative Fields
RequestMessage	service name, request ID, requester identity, payload, UserToken
RequestAckMessage	service name, request ID, provider identity, status, message, metadata payload, echoed UserToken, ProviderToken
SelectionMessage	service name, request ID, selected provider identity, ProviderToken
ResponseMessage	service name, request ID, provider identity, result payload, echoed UserToken

The message fields are intentionally service-independent. Application-specific data remains in the payload or in typed request/response objects, while the framework uses service names, request IDs, identities, and one-time tokens to provide authorization checks, replay resistance, and response verification.

3.6 Service Invocation Workflow

The normal NDNSF invocation workflow is:

1. The user publishes a service request under a service name.
2. Providers with permission to serve the service receive the request.
3. Providers decide whether to ACK based on availability, local state, and application-defined metadata.
4. The user collects ACKs during an ACK window.
5. A selection strategy chooses one or more providers.
6. Selected providers execute the request and publish responses.
7. The user verifies tokens, provider identity, and response data.

RequestNonce is the User-generated application-level random value; ProviderChallenge is the Provider-generated challenge checked through Selection. Their protected exchange is explained below. The table describes the single-Provider invocation. Plan identity, role assignments, and dependency checks are added by the collaboration extension below, not by the initial four-message flow.

These are NDNSF Data objects. Sync supports discovery of their names, after which Interests retrieve Data, possibly from a cache. A negative ACK is an application participation decision, not a network Nack. Not every observed Request must receive a reply: invalid, unauthorized, duplicate, or expired messages are subject to validation and state checks. A positive ACK expresses willingness; it does not itself start the selected task or guarantee later resource availability.

3.4 Authorization and Confidentiality

3.4.1 Rationale for ABE-Backed Service Authorization

NDNSF uses ABE to protect discovery descriptors and reuses attribute-based service permissions in its invocation checks. The first choice addresses confidentiality before the participating Providers are known; the second keeps service entitlements independent of identity-signing credentials. Neither implies that ABE is always preferable to role-based authorization with suitable encryption. Withdrawal and rekeying are lifecycle costs to evaluate. Execution authorization determines whether a User may invoke a service; Provider authorization determines whether an entity may offer that service. Content confidentiality limits who can read protected content. The Controller administers these service permissions; the User and Provider check them during invocation. Signature validation under trust rules and invocation-state checks remain necessary. Permission to offer a service is distinct from selection for one particular task.

The design addresses three requirements: keep discovery descriptors non-public before participating Providers are identified; authenticate participants, verify their service permissions, and reject replayed invocation messages; and support permission grants and withdrawals during

This workflow supports both one-to-one and one-to-many service models. The same request can reach multiple providers, enabling the user to select the first available provider, randomly balance load, select all providers, or use a custom strategy based on metadata.

The implemented provider-selection strategies are deliberately generic. The *First Responding* strategy selects the first successful ACK, which is useful when low latency is the main goal. *Random Selection* chooses among valid ACK candidates, which can spread load without requiring detailed provider state. *All Selected* chooses every provider that returned a valid ACK, which supports cooperative or redundant execution. A *Custom Selection* strategy runs after the ACK collection window and can use ACK metadata such as latency, queue length, load, location, model availability, safety state, memory, or accelerator capacity. The key design point is that readiness reporting and final provider choice are separate framework steps: providers decide whether and how to ACK, and the user-side strategy decides which ACK candidates should be selected.

3.7 Security Model

NDNSF uses a controller-authorized permission model. The ServiceController issues encrypted permission responses to target identities. Users and providers fetch their respective permissions, and permission Data is signed by the controller. NAC-ABE attributes map service messages to service-level authorization namespaces. One-time user and provider tokens protect the Request/ACK/Selection/Response flow against replay and confused-deputy behavior.

NAC-ABE enforcement is tied to the invocation tokens rather than to a separate permission list carried in every service message. A Request message includes a fresh `UserToken` encrypted under the service attribute, so only an authorized provider can decrypt the request and return that token. The ACK echoes the `UserToken` and carries a fresh `ProviderToken` encrypted under the permission attribute. The user proves selection authority by returning the selected provider’s `ProviderToken` in the Selection message, and the final Response echoes the `UserToken`. In this way, NAC-ABE attributes gate who can decrypt

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

This workflow supports both one-to-one and one-to-many service models. The same request can reach multiple providers, enabling the user to select the first available provider, randomly balance load, select all providers, or use a custom strategy based on metadata.

The implemented provider-selection strategies are deliberately generic. The *First Responding* strategy selects the first successful ACK, which is useful when low latency is the main goal. *Random Selection* chooses among valid ACK candidates, which can spread load without requiring detailed provider state. *All Selected* chooses every provider that returned a valid ACK, which supports cooperative or redundant execution. A *Custom Selection* strategy runs after the ACK collection window and can use ACK metadata such as latency, queue length, load, location, model availability, safety state, memory, or accelerator capacity. The key design point is that readiness reporting and final provider choice are separate framework steps: providers decide whether and how to ACK, and the user-side strategy decides which ACK candidates should be selected.

3.7 Security Model

NDNSF uses a controller-authorized permission model. The ServiceController issues encrypted permission responses to target identities. Users and providers fetch their respective permissions, and permission Data is signed by the controller. NAC-ABE attributes map service messages to service-level authorization namespaces. One-time user and provider tokens protect the Request/ACK/Selection/Response flow against replay and confused-deputy behavior.

NAC-ABE enforcement is tied to the invocation tokens rather than to a separate permission list carried in every service message. A Request message includes a fresh `UserToken` encrypted under the service attribute, so only an authorized provider can decrypt the request and return that token. The ACK echoes the `UserToken` and carries a fresh `ProviderToken` encrypted under the permission attribute. The user proves selection authority by returning the selected provider’s `ProviderToken` in the Selection message, and the final Response echoes the `UserToken`. In this way, NAC-ABE attributes gate who can decrypt

invocation timeout and result callbacks, while NDNSF manages the surrounding exchange. A service requiring confidential execution input uses the protected input-delivery path described below; this compact example does not prescribe placing arbitrary application input in discovery. Typed serialization removes some repeated application glue, but does not establish semantic compatibility between different service versions. That remains part of the service contract.

3.2 Participants, Trust, and Exclusions

These are protocol roles, not necessarily separate devices: one node may be both a User and a Provider. The managed organization registers identities and approves service permissions; the Controller administers that policy. Participants require a provisioned trust anchor and validation rules for identity certificates and Controller material. Not knowing which candidates will respond to this invocation does not mean that the authorization authority has never enrolled those entities. Billing, arbitrary business transactions, globally optimal scheduling, and protection against all denial-of-service traffic are outside the claimed scope. The User requests a service and selects its Provider; in the collaboration extension, it also coordinates role assignments. Providers offer and execute services. A Controller administers authorization material and signed status. Forwarders and caches can deliver Data without being trusted to authorize execution. The model trusts the Controller and the User’s planning decision; it does not claim consensus against a malicious coordinator. An attacker may observe visible names and traffic, inject invalid objects, replay previously valid messages, or operate without the needed service permission. Credential compromise and intentional sharing by an authorized principal require additional controls and are not prevented by ABE alone.

Authorization is permission to perform an operation, authentication establishes the signing identity under trust rules, and confidentiality limits who can recover Content. Valid credentials do not prove honest load reports, correct inference, or truthful sensor observations. Signed status must be obtained and validated before it can affect local authorization; a disconnected participant cannot be assumed to learn revocation immediately. Deadline/expiry

each authorization token, while one-time tokens bind that authorization check to a single service invocation.

Figure 4 illustrates the intended policy model: provider permissions specify which named services an entity may offer, while user permissions specify which named services an entity may access. This proposal keeps that service-name-centered policy model in the UAV and distributed-inference extensions, while adding deployment procedures that make the trust root and publisher namespaces explicit.




Figure 4: Example provider and user access-control policies in NDN SF.

The current model separates:

RequestNonce is the User-generated application-level random value; ProviderChallenge is the Provider-generated challenge checked through Selection. Their protected exchange is explained below. The table describes the single-Provider invocation. Plan identity, role assignments, and dependency checks are added by the collaboration extension below, not by the initial four-message flow.

These are NDN SF Data objects. Sync supports discovery of their names, after which Interests retrieve Data, possibly from a cache. A negative ACK is an application participation decision, not a network Nack. Not every observed Request must receive a reply: invalid, unauthorized, duplicate, or expired messages are subject to validation and state checks. A positive ACK expresses willingness; it does not itself start the selected task or guarantee later resource availability.

3.4 Authorization and Confidentiality

3.4.1 Rationale for ABE-Backed Service Authorization

NDN SF uses ABE to protect discovery descriptors and reuses attribute-based service permissions in its invocation checks. The first choice addresses confidentiality before the participating Providers are known; the second keeps service entitlements independent of identity-signing credentials. Neither implies that ABE is always preferable to role-based authorization with suitable encryption. Withdrawal and rekeying are lifecycle costs to evaluate. Execution authorization determines whether a User may invoke a service; Provider authorization determines whether an entity may offer that service. Content confidentiality limits who can read protected content. The Controller administers these service permissions; the User and Provider check them during invocation. Signature validation under trust rules and invocation-state checks remain necessary. Permission to offer a service is distinct from selection for one particular task.

The design addresses three requirements: keep discovery descriptors non-public before participating Providers are identified; authenticate participants, verify their service permissions, and reject replayed invocation messages; and support permission grants and withdrawals during

- **user authorization:** whether the user may request a service;
- **provider authorization:** whether a provider may serve a service;
- **message integrity:** whether a service message is signed and valid;
- **data confidentiality:** whether the payload is encrypted for authorized participants;
- **replay protection:** whether one-time tokens have already been consumed.

3.8 Completed Foundation: NDN v0.1 Evaluation

The completed NDN evaluation compares NDN against NSC and gRPC in MiniNDN. The evaluation focuses on service latency and success rate under wired, lossy, intermittent-provider, wireless mobility, adaptive-admission, selective ACK, and provider-selection scenarios. The key result is that NDN can outperform NSC and improve reliability under provider diversity and intermittent availability, while maintaining reasonable latency relative to IP-based RPC in lossless conditions.

The strongest evidence comes from scenarios where endpoint binding or a single provider becomes fragile. In the lossy wired Memphis–CSU–UCLA path, packet round-trip loss is $1 - (1 - \ell)^4$ for per-link loss ℓ . NDN completed 599 of 600 requests at 1% link loss, where the packet round-trip loss is 3.94%, and 598 of 600 requests at 3% link loss, where the packet round-trip loss is 11.47%. At 10% link loss, where packet round-trip loss rises to 34.39%, the best measured NDN result was 539 of 600 requests. This identifies the bounded-time delivery limit of SVS-based service invocation under heavy loss, rather than a failure of the NDN service model. In a wired multi-provider experiment with intermittent provider rejection, NDN maintained 99.46% success at 20% provider-failure probability and 96.70% success at 30%, whereas the single-provider baselines dropped to roughly two-thirds success or lower. In the wireless mobility experiment, NDN achieved 96.25% success at 100 m AP range and 100% at larger ranges, while NSC and gRPC were much lower when bound to one mobile provider.

operation. Permission aggregation is an additional management benefit, not a fourth security guarantee.

Reason 1: confidential discovery before Provider selection. Before Provider selection, the User publishes a discovery descriptor containing the location, deadline, and resource requirements needed for a Provider’s positive-ACK decision. For example, nearby UAVs need an observation area and arrival deadline to decide whether to participate, but do not yet need the full application input. The User encrypts this descriptor using ABE and the requested service’s attribute, without enumerating Provider identities. Confidentiality is against parties lacking the required decryption capability: authorized but unselected Providers can still read the descriptor. Content encryption does not hide outer Data names or traffic patterns. In the request-scoped confidentiality path, the full application input is encrypted with a request-scoped key; Selection protects that key for the selected Provider. The requesting User retains access. A Trust Schema alone does not encrypt content; a role-authorized design can also use ABE or service-group encryption to provide discovery confidentiality.

Encrypted Request Data can be retrieved or cached without granting decryption access. This separation of delivery from disclosure also holds for other content-encryption schemes; it is not unique to ABE.

Reason 2: reusing attribute-based permissions for service invocation. Confidential discovery alone does not justify ABE-backed execution authorization: ABE encryption could instead be combined with role-based invocation checks. NDN chooses to administer both Provider and User service permissions through the existing attribute-authority infrastructure, separately from identity-signing credentials.

Provider permissions determine access to RequestNonce; User permissions determine access to ProviderChallenge. NDN checks RequestNonce in ACK Data and validates Selection Data using the stored ProviderChallenge. Signature, current-permission, and request-state checks remain separate requirements (Section 3.4.2). Decryption is therefore one part of the invocation

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

check, not permission to execute any request.

DNMP instead uses role signing keys and a Trust Schema to constrain command targets and actions [10]. In the DNMP-inspired alternative (Section 2.4.2), retaining ABE for discovery would require managing both decryption permissions and role-based invocation permissions. NDNFSF uses the attribute-authority infrastructure for both service-level checks, without adding role signing credentials solely to represent service membership. Trust rules for validating Data signatures remain necessary.

This is a management-reuse rationale, not a demonstrated cost advantage. Provider and User permissions remain distinct; discovery encryption does not already supply User permission decryption keys or challenge processing. For policies requiring command-argument restrictions, DNMP’s schema-based checks may be a better fit than NDNFSF’s service-level checks. The evaluation must compare both designs at the same policy granularity, including credential updates and invocation overhead.

Reason 3: service-permission aggregation as an additional management benefit. KP-ABE associates attributes with ciphertext and an access policy with each decryption key (DKEY) [19]. NDNFSF uses /SERVICE/<service> for Provider access and /PERMISSION/<service> for User access. Within one authority and parameter generation, an entity’s service permissions can be combined in one individually issued DKEY. For example, one User’s policy can allow detection OR mapping, while another allows detection only. Their identity-signing credentials need not change to represent these different service sets; each service retains its own ciphertext attribute.

Compared with DNMP’s role-key-based authorization [10], NDNFSF represents service entitlements in the DKEY rather than in additional role-specific signing credentials. Where equivalent entitlements require several distinct role credentials, this avoids issuing and maintaining those additional signing credentials. It does not imply one credential per DNMP service: one role can permit several commands. Nor does it establish a universal advantage over

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Trust Schemas, which can express other credential arrangements. A single DKEY does not have fixed size or update cost; policy complexity, distribution, and rekeying still matter. The benefit concerns independently managed credentials, not a demonstrated reduction in total key bytes or runtime cost.

3.4.2 Authorization Mechanisms, Lifecycle, and Costs

Identity authentication. The User and Provider validate Data signatures and certificate chains under the applicable trust rules. This identifies an accepted signer; it does not alone establish that signer’s permission to offer or invoke the requested service.

Service-permission verification. In the ABE-backed invocation path, the User generates a fresh random value, *RequestNonce*, and includes it in ABE-protected Request content. This is an NDNSF application-level value, not the NDN Interest Nonce. A Provider that recovers it returns it in ACK Data; the User compares the returned value with the one stored for the active request. Thus RequestNonce supports both request binding and a check of access to the service-protected value. The Provider generates *ProviderChallenge* and includes it in ABE-protected ACK content. The User uses the recovered ProviderChallenge to construct Selection Data. Before execution, the Provider validates Selection Data using ProviderChallenge stored for that request. Together with current service-permission checks, these checks determine whether the Provider may offer the service and whether the User may invoke it.

These checks demonstrate access to protected challenge values, not the creation of new permissions or a formal proof of exclusive key possession. Authorized key or challenge sharing is outside the claimed threat protection. This reuses attribute-based permission management rather than a separate set of role-based invocation rules. It is not cost-free: confidential discovery already needs Provider DKEYs, whereas User permission DKEYs and challenge processing are additional authorization costs. ABE confidentiality combined with DNMP-inspired role authorization remains a valid alternative.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Replay protection. Participants match messages to an active request and its protocol state. The Provider records whether ProviderChallenge has already been used to accept Selection Data for execution; receiving the same Selection again must not start another execution. This requires appropriate retention and invalidation of request state, including across restart. Neither encryption nor possession of a challenge value alone prevents replay. These admission checks do not establish exactly-once application side effects across crashes.

Runtime permission grants and withdrawals. The Controller updates permission status for both operations. An ordinary grant updates the target entity’s complete permission policy; that entity retrieves a replacement DKEY without changing the ABE public parameters. Withdrawal additionally rotates the Controller’s ABE master key and public parameters. Entities retaining permissions obtain replacement DKEYs encoding their remaining permissions. The master key stays with the Controller. Rotation is Controller-wide, not confined to the withdrawn service.

In the revised dissemination design, the Controller publishes updated public parameters as signed named Data and uses Sync to announce their availability. Participants retrieve and validate the parameters and Controller-signed PolicyStatusData, then refresh their DKEYs as required. Ordinary grants also require participants to discover and retrieve updated permission status, even though the public parameters have not changed. Sync is an update-discovery mechanism, not the authority for granting permission. The local implementation supports status retrieval and parameter rotation; the dedicated Sync announcement path remains to be verified end to end.

Current-permission checks enforce withdrawal after the updated status has been retrieved and validated. Parameter rotation restricts exchanges encrypted under the new generation. Disconnected nodes cannot apply updates immediately, and previously disclosed keys or plaintext cannot be taken back.

Lifecycle comparison and evaluation. DNMP specifies validity bounds and replay-state handling [10]; these do not by themselves provide immediate withdrawal of an unexpired role

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

key. For the comparison, a DNMP-inspired design can revoke the relevant credential and require execution nodes to check signed revocation status or reject expired certificates. This invalidates authority at the validator; it does not physically retrieve distributed certificates. Both approaches require update propagation, and neither retracts previously disclosed keys or plaintext.

Focused integration tests exercise NDNSF’s runtime revocation checks, including identity-, certificate-, and service-scoped withdrawal. These tests establish behavior at the exercised runtime boundaries, not network-wide instantaneous revocation or lower cost than certificate-based alternatives. The remaining evaluation should measure status-propagation delay, rejection delay, rekeying traffic, and effects on unaffected members. Runtime status validation and key rotation support runtime permission changes, with explicit propagation and rekeying costs. They are not ABE-exclusive capabilities. NDNSF introduces ABE for confidential discovery and reuses its service-permission management for invocation checks and permission updates; signatures and request-state checks remain necessary, separate parts of the design.

Discovery confidentiality applies to descriptor Content against parties lacking the required ABE decryption capability. It does not conceal the descriptor from other authorized candidates or provide traffic-analysis resistance. In the request-scoped protection design, the User encrypts application input with a symmetric key and delivers the necessary key material through Selection, encrypted for the chosen Provider’s authenticated public key. The response key is shared by the requesting User and selected Provider. Other Users cannot decrypt the protected result under the key-secrecy assumptions; the Provider necessarily knows the plaintext it produces. These are the protection requirements for the evaluated design, rather than claims about every historical implementation path.

3.5 Invocation Modes and Execution Lifecycle

3.5.1 Discovery, Known Providers, and Local Composition

Normal discovery is appropriate when the executor has not been chosen. The User collects ACKs within a bounded window and chooses among positive ACKs that pass the

- **Provider collaboration:** v0.1 supports selecting providers, but larger tasks need collaboration among multiple providers, partial failure handling, and compensation requests.
- **Service-container composition:** real applications often contain multiple services in one trusted process. Those services need efficient local composition without exposing a network bypass to remote callers.
- **Large artifacts:** v0.1 service payloads are not intended to carry models, recordings, images, or intermediate tensors. Applications need typed references to named Data objects plus native NDN segmented retrieval for large signed and encrypted objects.
- **Operational validation:** v0.1 demonstrates the security model, but broader validation needs clearer certificate bootstrapping, trust-schema guidance, configuration procedures, and reproducible experiments.
- **Application-level validation:** v0.1 evaluation shows the core framework works in MiniNDN, but the dissertation must evaluate whether the design supports demanding UAV and distributed-inference workflows.

3.10 Supporting Mechanisms for Provider Collaboration

The following mechanisms support provider collaboration in specific operating conditions. They are motivated by the UAV and distributed-inference workloads, but they are not separate dissertation threads and should remain generic framework mechanisms rather than application-specific shortcuts.

3.10.1 Targeted Invocation

Some services, such as UAV flight-control commands, are directed at a known provider and must minimize latency. The proposed Targeted invocation path keeps NDNsf permissions,

key. For the comparison, a DNMP-inspired design can revoke the relevant credential and require execution nodes to check signed revocation status or reject expired certificates. This invalidates authority at the validator; it does not physically retrieve distributed certificates. Both approaches require update propagation, and neither retracts previously disclosed keys or plaintext.

Focused integration tests exercise NDNsf’s runtime revocation checks, including identity-, certificate-, and service-scoped withdrawal. These tests establish behavior at the exercised runtime boundaries, not network-wide instantaneous revocation or lower cost than certificate-based alternatives. The remaining evaluation should measure status-propagation delay, rejection delay, rekeying traffic, and effects on unaffected members. Runtime status validation and key rotation support runtime permission changes, with explicit propagation and rekeying costs. They are not ABE-exclusive capabilities. NDNsf introduces ABE for confidential discovery and reuses its service-permission management for invocation checks and permission updates; signatures and request-state checks remain necessary, separate parts of the design.

Discovery confidentiality applies to descriptor Content against parties lacking the required ABE decryption capability. It does not conceal the descriptor from other authorized candidates or provide traffic-analysis resistance. In the request-scoped protection design, the User encrypts application input with a symmetric key and delivers the necessary key material through Selection, encrypted for the chosen Provider’s authenticated public key. The response key is shared by the requesting User and selected Provider. Other Users cannot decrypt the protected result under the key-secrecy assumptions; the Provider necessarily knows the plaintext it produces. These are the protection requirements for the evaluated design, rather than claims about every historical implementation path.

3.5 Invocation Modes and Execution Lifecycle

3.5.1 Discovery, Known Providers, and Local Composition

Normal discovery is appropriate when the executor has not been chosen. The User collects ACKs within a bounded window and chooses among positive ACKs that pass the

one-time user/provider tokens, and replay protection, but avoids repeated ACK/Selection rounds after a token batch has been bootstrapped through the normal flow. This gives known-provider commands a fast path while keeping the security semantics of normal NDN SF invocation.

3.10.2 *Trusted Local Invocation*

A service-container application may contain several services in the same trusted process. A local call from one internal service component to another should not need to publish an NDN request. The proposed local invocation path therefore provides synchronous and asynchronous in-process invocation for explicitly registered local services. It is not a wire-protocol mode: remote callers cannot request it, and it bypasses NDN SF network security only inside the trusted local container boundary.

3.10.3 *Large-Data Reference Abstraction*

Models, video, images, activations, and logs should not be placed in small request payloads. The proposed pattern is to let the producer segment the object as named NDN Data and place only a typed object reference in the NDN SF request or response. The reference can include the object name, object type, version or session id, expected digest, expected size, final segment information, source type, and optional encryption metadata. The receiver then uses normal NDN segmented fetching, such as ndn-cxx segment fetch logic, to retrieve the object by name.

This abstraction is intended to be a generic NDN SF mechanism, not an API for a particular storage service. Small request and response payloads remain useful for control messages and compact typed objects. When the payload becomes a large artifact, the same logical service invocation should carry a *reference* rather than the artifact bytes. The application or helper can turn a payload into named Data segments, sign and optionally encrypt those segments, and then pass a typed reference through NDN SF. This keeps service selection, authorization, token checks, and timeout diagnosis in the normal NDN SF path, while large bytes move through NDN’s native named-data retrieval path.

applicable validation and policy checks. A negative ACK communicates non-participation; silence cannot distinguish unwillingness, loss, and disconnection. A selection strategy chooses among candidates for the requested operation. Choosing several independent executions, where supported, is distinct from assigning complementary roles with dependencies in a collaboration plan.

Known-Provider invocation is appropriate when application meaning fixes the device, for example, stopping the camera that was previously started. The current `RequestServiceTargeted` path can reuse previously bootstrapped one-time authorization material to avoid repeated ACK/Selection exchanges under its supported conditions. It still checks current authorization and request state. The fast path is not universal: request-scoped confidentiality needs a Selection envelope to carry the input and response keys, and the current implementation uses the bootstrap/ACK/Selection path for that case. The explicitly Selection-gated input mode is rejected by the Selection-free Targeted interface. Thus, targeting a known Provider is an application requirement, whereas omitting exchanges is a conditional implementation optimization.

Trusted local composition operates within a process through the local registry. It has neither remote discovery nor the same network threat boundary. Evaluation should identify whether a handler was reached locally, by normal discovery, or by a targeted exchange. Otherwise an apparent invocation improvement might merely measure a different operation.

3.5.2 *Admission, Progress, and Failure*

An API submission, admission to a runtime queue, acceptance by a Provider, and completion of application work are different events. Under overload, reporting latency only for completed requests can hide dropped or rejected work. The runtime therefore needs observable outcomes for refusal, timeout, and terminal results, while an evaluation must count all submitted requests in its stated denominator. Selection-status queries can help determine the observed status of selected work; they are not proof of a physical effect or a distributed transaction.

The same reference format should be usable across UAV recordings, camera artifacts, model shards, runtime artifacts, intermediate activations, and logs. As a result, UAV-APP and DistributedInference can share one framework-level large-data vocabulary: object type, object name, producer, session or version, size, digest, final segment hint, and optional encryption metadata. This avoids embedding application-specific conventions such as raw Data-name strings inside service payloads and makes large-object behavior inspectable by applications that do not understand NDN packet details.

This design can optionally use persistent storage, but storage is an implementation support component rather than a separate research direction in this proposal. If persistent storage is used, it stores opaque NDN Data segments that have already been named, signed, and optionally encrypted by the application or by a framework helper. Security is not handled by a storage- specific path: NDN SF’s normal hybrid encryption, signatures, tokens, permissions, and trust-schema validation protect the service message and the referenced Data. Discovery metadata can still accelerate validation, but large- object transfer itself is based on typed object references and native NDN segmented retrieval.

This chapter supports RQ2 by previewing the reusable NDN SF mechanisms that the dissertation will refine: secure invocation, provider authorization, provider collaboration, and its supporting mechanisms. It supports RQ3 by identifying where collaboration introduces reliability, security, and latency tradeoffs, and where supporting mechanisms are needed to control overhead without weakening the core service semantics.

Applications must also distinguish cancellation of interest in a result from reversal of an action. A recording can be stopped, but previously recorded frames still exist. An inference worker can stop producing new output, but a completed remote action cannot be undone by deleting requester state. Retrying requires the application’s stated idempotency or compensation policy. For the proposed collaboration extension, replacement additionally requires old assignments and late callbacks to stop advancing the current plan.

3.6 Control Messages, Streams, and Complete Objects

Small service messages communicate intent, permission evidence, and status. Large application data should not be copied into each control message. An object reference allows a service to identify a model fragment, recording, or image, while the recipient retrieves its named Data separately. The reference and service contract must supply enough information to select the intended version and validate the reconstructed object. The concrete interface determines the available name, size, digest, segmentation, and protection metadata; not every object type needs an identical collection of fields.

The important separation is between identifying an object, obtaining its bytes, and authorizing its use. A digest can detect a mismatch with the expected bytes only when the expected digest itself is trusted. A signature binds Data to a key under the application’s trust rules. The invocation or collaboration checks establish whether that producer and object may satisfy the current input. Recipient-scoped key delivery controls access to protected Content. None of these checks should silently substitute for another.

An ongoing stream has a different lifecycle from a complete object. Live video may prefer recent frames and tolerate explicitly identified gaps. A model artifact generally requires every needed byte before loading. Incremental inference requires ordered application events and a terminal outcome, but its token stream is not necessarily interchangeable with a camera stream interface. NDN SF supplies named-data mechanisms; applications define acceptable gaps, buffering limits, and completeness rules. The evaluation must measure the chosen policy,

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

including its discarded data and queueing costs.

Named-object reuse is valuable in the multi-view case: the compute Provider returns annotations and references to the original images, with authorized access to their content keys. The requester can retrieve those images from available producers or caches and validate their original signatures. This avoids producing a second image object merely to change the receiver. It does not guarantee that a cached copy remains available; retention requirements need an explicit producer or repository policy.

3.7 Proposed Extension: From ACKs to a Collaboration Plan

The extension changes the selection problem from choosing one executor to assigning complementary task roles. The application supplies a role/dependency template, required services, feasibility constraints, and deadline. The template may be written manually; automatic generation is not necessary to establish the framework contribution. A replaceable planning policy selects among eligible positive ACKs after the collection window closes. Closing the window freezes the candidates observed for that decision, not all possible Providers in the network. ACK claims can become stale, so a selected Provider still performs execution-time admission and may reject work it can no longer execute.

In the proposed design, the User signs a Plan Manifest containing the request identity, accepted candidate snapshot, exact role-to-Provider assignments, permitted named dependency edges, authorization/key scope, and execution identity. Its digest is computed over a canonical plan body that excludes the digest field itself. Providers validate that signed plan before deriving local dependency checks; a long-term certificate does not specify this task’s assignments. The User must record the committed plan and reject a different plan for the same execution identity. This is a local state-machine rule under the trusted-User model, not a consequence of signing alone or protection against a Byzantine User signing two plans.

3.3 Provider Collaboration as the Central Extension

The central extension beyond NDN v0.1 is provider collaboration. In a collaborative service, one logical request may involve multiple authorized providers rather than a single endpoint. Providers may advertise ACK metadata such as availability, queue state, location, or role capability; the user may select one or more providers after an ACK collection interval; and the application may issue compensation requests when only part of a logical task is completed. Long-running selected work should also be diagnosable: after a timeout, the user should be able to learn whether a selected provider has not seen the selection, is queued, is executing, has failed, or has already published a response.

This collaboration model is intentionally application-independent. Multi-UAV patrol can use it to assign mission segments and recover from partial failure, while distributed inference can use it to assign model roles and coordinate intermediate artifacts. In both cases, NDN should provide the generic service semantics: request publication, provider metadata, provider selection, authorization, status diagnosis, and response verification. The application defines the domain-specific work; the framework defines how selected providers participate in one logical service workflow.

3.4 Generic Dynamic API

NDN v0.1 moved toward a generic dynamic API instead of generated, service-specific stubs. A provider can register a typed service handler:

```
provider.addHandler<RequestT, ResponseT>(
    ndn::Name("/ObjectDetection/YOLOv8"),
    handler);
```

A user can invoke a service through the generic runtime:

```
user.RequestService<RequestT, ResponseT>(
    providers,
```

including its discarded data and queueing costs.

Named-object reuse is valuable in the multi-view case: the compute Provider returns annotations and references to the original images, with authorized access to their content keys. The requester can retrieve those images from available producers or caches and validate their original signatures. This avoids producing a second image object merely to change the receiver. It does not guarantee that a cached copy remains available; retention requirements need an explicit producer or repository policy.

3.7 Proposed Extension: From ACKs to a Collaboration Plan

The extension changes the selection problem from choosing one executor to assigning complementary task roles. The application supplies a role/dependency template, required services, feasibility constraints, and deadline. The template may be written manually; automatic generation is not necessary to establish the framework contribution. A replaceable planning policy selects among eligible positive ACKs after the collection window closes. Closing the window freezes the candidates observed for that decision, not all possible Providers in the network. ACK claims can become stale, so a selected Provider still performs execution-time admission and may reject work it can no longer execute.

In the proposed design, the User signs a Plan Manifest containing the request identity, accepted candidate snapshot, exact role-to-Provider assignments, permitted named dependency edges, authorization/key scope, and execution identity. Its digest is computed over a canonical plan body that excludes the digest field itself. Providers validate that signed plan before deriving local dependency checks; a long-term certificate does not specify this task’s assignments. The User must record the committed plan and reject a different plan for the same execution identity. This is a local state-machine rule under the trusted-User model, not a consequence of signing alone or protection against a Byzantine User signing two plans.

Table 6: Selected mechanism results from the submitted NDNNSF v0.1 evaluation.

Mechanism	Baseline Behavior	NDN NSF Behavior
Adaptive admission under 100 RPS offered load	Without admission control, P95 latency reached 4254.39 ms and timeout rate reached 3.55%.	With admission control, admitted rate was limited to about 81 RPS, P95 latency stayed at 268.12 ms, and timeout rate was 0.00%.
Selective ACK with nearest provider overloaded	ACK-all selected the nearest overloaded provider for 2209 of 2250 requests and achieved 1.82% success.	Selective ACK suppressed overloaded ACKs, shifted selection to available providers, and achieved 100.00% success.
Real-machine intermittent availability	Single-provider NSC/gRPC stayed near 50% success when the chosen provider was unavailable for half the run.	NDN NSF selected among two providers and reached 74.62% success at 10 RPS and 73.73% at 30 RPS, close to the expected availability window.

3.9 Limitations of NDNNSF v0.1

NDNNSF v0.1 establishes feasibility, but it does not yet fully address the requirements of dynamic service applications beyond controlled framework tests. The main limitations that motivate the remaining dissertation work are:

- **Low-latency known-provider control:** the normal Request/ACK/Selection/Response workflow is appropriate for provider discovery, but control commands to a known drone or device need a faster Targeted path while preserving authorization and replay protection.

3.8 Proposed Dependency Admission and Execution State

Under the proposed acceptance rules, a downstream Provider checks that the execution is active and within its deadline before accepting a dependency. The Data signature must validate under the applicable trust rules; the request, plan, and execution identity must match the assignment; the signer must be the assigned predecessor; and the name, data format, and key scope must match a declared edge. Required segments must validate, and incomplete or malformed objects cannot satisfy an edge. The same checks apply to producer-, cache-, and repository-fetched Data. The plan references objects in their actual producers’ authorized namespaces rather than republishing them under a different producer identity.

Table 4: Proposed task-state contract and required negative checks.

State/event	Permitted transition	Required rejection or bounded behavior
Collecting ACKs	Close window and validate candidate set	Late ACK does not silently alter a committed plan
Selected	Validate Selection and perform local admission	Positive ACK alone, wrong Provider, or spent challenge cannot start work
Waiting for dependencies	Execute when required authorized edges are satisfied	Wrong role, undeclared edge, incomplete object, or wrong active execution is rejected
Completed/cancelled	Return recorded status under duplicate-message policy	Duplicate publication does not implicitly repeat a side effect
Failure/reassignment	Explicitly terminate or authorize a replacement execution	Old execution must not acquire authority from the new assignment; allowed result reuse is explicit

These are proposed acceptance conditions, not a claim that every production path is qualified. Network re-retrieval is not identical to repeating computation. If a task is reassigned, cancellation, fencing, and result-reuse rules must be explicit. An already computed image may remain useful when the new plan explicitly authorizes its reuse; an old object is not automatically invalid solely because it is cached. Application-defined compensation can limit side effects but does not provide distributed atomic transactions or exactly-once physical action.

- **Provider collaboration:** v0.1 supports selecting providers, but larger tasks need collaboration among multiple providers, partial failure handling, and compensation requests.
- **Service-container composition:** real applications often contain multiple services in one trusted process. Those services need efficient local composition without exposing a network bypass to remote callers.
- **Large artifacts:** v0.1 service payloads are not intended to carry models, recordings, images, or intermediate tensors. Applications need typed references to named Data objects plus native NDN segmented retrieval for large signed and encrypted objects.
- **Operational validation:** v0.1 demonstrates the security model, but broader validation needs clearer certificate bootstrapping, trust-schema guidance, configuration procedures, and reproducible experiments.
- **Application-level validation:** v0.1 evaluation shows the core framework works in MiniNDN, but the dissertation must evaluate whether the design supports demanding UAV and distributed-inference workflows.

3.10 Supporting Mechanisms for Provider Collaboration

The following mechanisms support provider collaboration in specific operating conditions. They are motivated by the UAV and distributed-inference workloads, but they are not separate dissertation threads and should remain generic framework mechanisms rather than application-specific shortcuts.

3.10.1 Targeted Invocation

Some services, such as UAV flight-control commands, are directed at a known provider and must minimize latency. The proposed Targeted invocation path keeps NDNsf permissions,

The authorization design and comparative evaluation are intended to answer RQ1; the proposed plan/state rules and their tests are intended to answer RQ2. Neither question is answered by a design description alone. Existing invocation evidence does not establish collaboration correctness; wrong-role, wrong-edge, replay, cancellation, and direct/cache tests remain separate evaluation obligations.

3.9 Conditional Correctness Argument

Assume uncompromised signing keys, correctly enforced trust/authorization checks, and Providers that obey the specified state machine. Selection-before-execution follows only if every handler-start path checks an active assignment and unspent authorization state. Dependency integrity follows only if every input path, including a cache hit, checks the assigned producer and declared edge before satisfying the dependency. Replacement isolation additionally requires old workers and late callbacks to lose authority under the new execution identity. Restart recovery must preserve or safely invalidate the relevant one-time state; signature validity and cache FreshnessPeriod alone are insufficient.

These are conditional invariants and implementation proof obligations, not a cryptographic proof of a complete system. Tests should inject a wrong-role object within the same request, an undeclared edge from a correctly signed producer, and a late callback after replacement. Each isolates a property that a request ID alone cannot express. Application result correctness and physical side effects remain outside this argument.

Chapter 4

UAV Network Application over NDN

4.1 Why UAV Networks Need a Service Framework

UAV applications are a strong test case for NDN because they combine mobility, real-time control, video streaming, distributed sensing, multi-agent coordination, and safety constraints. A UAV network may include a ground station, several drones, cameras, flight controllers, object-detection services, storage capabilities, and mission-planning components. Each component exposes capabilities that naturally look like services.

For the dissertation, the UAV workload is specifically used to stress provider collaboration under mobile and command-sensitive conditions. A ground station may need to discover multiple drones, select one or more providers for a mission, issue targeted commands to known drones, compensate for a failed mission segment, and retrieve video or recording objects without treating large data as RPC payload. This makes UAV-APP a validation workload for the NDN framework: it asks whether provider collaboration, authorization, selection status, and supporting low-latency or large-data mechanisms can work together in an operator-facing service container.

Developing such an application over IP creates several recurring problems:

- **Address binding.** A drone’s identity and its IP address may change as it moves, reconnects, or switches networks.
- **NAT and mobility.** Direct connectivity to drones is not always available, especially in ad hoc or field deployments.
- **Service discovery.** A ground station needs to know which drones offer camera, telemetry, flight-control, and mission services.
- **Command authorization.** Flight-control commands must be restricted to authorized users and protected against replay.

Chapter 4

Application Scenarios and Validation Plan

4.1 Validation Strategy: Applications Exercise the Framework

We use NDN-UAV and NDN-DI to evaluate how NDN supports service discovery, authorization, data delivery, and execution management across both single-Provider and multi-Provider workflows. Single-Provider operations remain relevant to RQ1 and RQ3; workflows with complementary roles additionally examine RQ2. The multi-view and distributed-model cases below are representative workloads within these broader applications.

The evaluation maps each application function to the NDN path it actually uses and to an observable outcome. An application feature supplies evidence only when that path is exercised and the result is recorded. A local helper call, for example, does not establish remote invocation authorization. Likewise, a working interface or model does not by itself demonstrate framework correctness or performance. The chapter defines validation coverage; Chapter 5 distinguishes available measurements from remaining tests.

4.2 NDN-UAV: Control, Sensing, and Collaboration

4.2.1 Application Structure and Service Workflows

The UAV application connects an operator’s ground station with drone-side service containers. Drone services expose camera operations, telemetry, and mission-related operations through NDN. The ground station can also provide an object-detection service, so the same application includes both service consumption and service provision. Application code interprets camera, vehicle, and mission data; NDN supplies the reusable invocation and data-delivery mechanisms.

These operations cover different framework requirements. Selecting a currently available sensing Provider exercises discovery and ACK-based selection. Addressing an already chosen vehicle exercises invocation associated with that Provider rather than selecting an arbitrary

- **Data verification.** Telemetry, images, videos, and mission results should be verifiable independently of the channel.
- **Failure recovery.** If one drone disconnects or fails a task, the mission should compensate using other drones when possible.
- **Large data.** Live video and recorded video should use streaming and segmented data retrieval instead of small RPC payloads.

4.2 UAV Services

The proposed NDN-SF-UAV-APP models a UAV network as a service-container application. The ground station and each drone run containers with different service instances, and the APP-level configuration determines identities, service names, runtime roles, camera behavior, and flight-controller backend. Table 7 summarizes the service types.

Table 7: Service types in the proposed UAV application.

Service Type	Purpose
Flight control	Arm, takeoff, land, goto, manual control, and generic MAVLink execution.
Telemetry and state	Report GPS, altitude, battery, armed state, flight mode, readiness, camera availability, link state, and mission state.
Video control	Start/stop live video, negotiate bitrate, and report actual streaming state.
Video data	Publish encoded frames or chunks as NDN Data for low-latency playback.
Recording discovery	Authorized discovery of locally recorded encrypted video objects.
Recording retrieval	Retrieve encrypted signed NDN Data objects through large-data references.
Object detection	Detect objects such as cars and trucks from camera frames, either at the ground station or onboard.
Mission/patrol	Assign waypoints, partition tasks among drones, start/stop missions, and compensate for missing subtasks.
Discovery/status	Publish installed service information and readiness state.

substitute. Telemetry and mission-status retrieval exercise the association of returned Data with the intended vehicle and active operation. Long-running operations supply cases for cancellation, deadlines, and handling late status. Tests must distinguish rejection of an invalid command from the flight controller’s physical execution of an accepted command.

Live video and recordings exercise different data paths. Camera Start/Stop is a service-control operation; the resulting live stream uses the named-data streaming interface. Recording references and complete stored objects require exact object retrieval rather than assumptions about cache persistence. These functions therefore contribute to validating stream lifecycle, ordering and gap handling, protected content delivery, and data-product retrieval. Video codecs, image processing, operator interfaces, and flight safety remain application responsibilities; their use of NDN-SF is nevertheless part of framework validation.

4.2.2 Operator Workflow and State Interpretation

A typical session starts with an operator discovering sensing services and choosing a vehicle. Subsequent camera or mission operations retain that vehicle identity unless the operator or a stated mission policy authorizes replacement. The interface associates status with the vehicle and operation that produced it. A delayed observation must not be presented as evidence that a newly issued command succeeded. Observation time, operation identity, and the application’s stale-data threshold are therefore part of the validation contract.

The proposed evaluation follows a complete session rather than testing isolated buttons: obtain status, start a camera stream, inspect incremental frames, stop the stream, find a recording, and retrieve the selected recorded object. Camera control uses service invocation; frame delivery and recording retrieval use their respective data paths. An implementation may record while no live viewer is attached, so recording availability should not be inferred from stream subscription. The test records whether the configured application supports that mode and whether recording references identify the correct producer and version.

Flight-related services connect the application to the vehicle’s flight-control interface,

4.3 Architecture

The UAV application contains two main roles:

- **Ground station role:** operator interaction, map-based mission planning, telemetry and video observation, object-detection service composition, and NDNsf user behavior.
- **Drone role:** flight-control adaptation, telemetry publication, camera and video publication, encrypted recording support, mission execution, and NDNsf provider behavior.

The application is explicitly service-oriented. A drone does not expose one monolithic endpoint. Instead, it exposes independently named services for flight control, telemetry, video control, camera frames, recording discovery, and mission assignment. The ground station uses structured runtime state to represent telemetry freshness, mission progress, video availability, and safety readiness, so that operator-visible status and command decisions remain consistent across independently named services. Known-drone commands use Targeted invocation, while collaborative workflows such as patrol assignment use normal NDNsf selection.

4.4 Operator Workflow and State Discipline

The UAV application is not intended as a replacement for an existing ground control station. Its purpose is to validate whether NDNsf can support a real UAV service workload. For that validation to be meaningful, the GS must still answer the operator questions that matter in practice: which vehicle is selected, whether it is ready and safe to control, which command or mission is active, whether telemetry and video are fresh, and what will happen if a link or device fails. The proposed UAV work therefore treats the GS as a state-driven operator console, not only as a collection of service invocations.

Table 5: UAV workflow contracts to exercise in application validation.

Operation	Application meaning	Observable framework outcome
Telemetry and status	Observation of the selected vehicle	Correct producer and operation association; stale data identified
Camera Start/Stop	Change the selected camera’s activity	Authorized invocation and terminal status distinct from frame arrival
Live video	Recent frames from the requested stream	Start delay, useful frame delivery, gaps, bounded buffering
Recording retrieval	Obtain a particular stored capture	Exact object and producer validated; Content access enforced
Mission operation	Issue bounded work and follow progress	Admission, progress, timeout, and completion distinguished

such as MAVLink and PX4 [20, 21]. Authorization to submit a command is separate from its operational safety. Geofencing, vehicle readiness, collision avoidance, and local failsafe behavior belong to the vehicle/application safety design. The networking study should use bounded supervisory commands and an explicit loss-of-contact policy, rather than assume that a service invocation is suitable for every low-level flight-control loop.

Multi-vehicle patrol provides a secondary workload: an application divides a region into tasks, tracks their completion, and decides whether unfinished tasks can be reassigned. Its relevance is the association of task state and returned observations with the active assignment. This is broader than multi-view analysis, but it should not become an additional unbounded algorithm-development requirement. The primary collaboration experiment below fixes its task and analysis method so that network and protocol effects can be isolated.

For both workflows, failure experiments should include a missing observation, a disconnected selected UAV, and delayed status after a replacement decision. Report the time during which a feasible Provider is actually reachable as well as the request deadline. Retransmission and synchronization cannot manufacture radio coverage or extend an expired application time budget. A partial completion rate can therefore reflect availability limits even

Table 8: Operator-facing state discipline planned for NDN-SF-UAV-APP.

State Area	Planned Behavior
Vehicle selection	The GS keeps the operator’s selected vehicle stable while telemetry, video, and service updates arrive asynchronously.
Readiness	The GS presents an integrated readiness view for camera, flight controller, telemetry, mission, video, recording, and storage-related capabilities.
Telemetry freshness	The GS distinguishes current, stale, and missing telemetry so that old data is not mistaken for current state.
Command lifecycle	Control and mission commands expose observable progress, completion, and timeout outcomes instead of opaque success/failure messages.
Mission progress	Patrol missions expose assignment, progress, partial completion, cancellation, and compensation behavior across drones.
Failure behavior	Timeouts and device/link failures surface diagnostic information so that the operator or application policy can decide the next action.

4.5 Live Video and Recording

Live video is initiated by a GS command to a drone-specific video-control service. The drone negotiates the accepted bitrate and frame size, encodes camera frames as a streaming codec, and publishes stream Data under names that include a stream identifier and monotonically increasing packet sequence. Each packet also carries stream metadata so the receiver can reject stale packets from a previous stream. The GS fetches and decodes video with an adaptive low-latency policy that uses observed RTT, bitrate, timeout pressure, decoder backlog, and published/decoded gap to tune the prefetch window and stale-frame skipping policy.

Recording is intentionally different from live control. The drone may keep its camera

when retrieval recovery behaves correctly.

4.2.3 Representative Collaboration: Multi-View Collection and Analysis

An operator chooses a target location or area on a map. Request Data carries only the service-protected descriptor needed to assess coverage, viewing feasibility, sensors, and timing. Eligible UAVs return positive ACKs. The proposed workflow assigns complementary viewing roles and a processing role, then delivers execution details and appropriate keys through Selection.

Each sensing UAV signs encrypted ViewImage Data under its own authorized namespace, with observation time and available pose/calibration metadata. The compute Provider retrieves permitted images, performs the specified analysis, and publishes signed Annotation Data. Its encrypted result supplies original image references and permitted image keys to the requester, which retrieves the images and renders annotations. The workflow reuses the original signed images rather than republishing them under another producer’s identity.

This case examines role assignments, named dependencies, result reuse, and incomplete-view or processing failures. It does not replace the control, telemetry, video, and recording validation above. GPS alone does not establish camera visibility, depth, or correspondence between views. Improved recognition or localization requires ground truth, a single-view control, and a specified fusion method. MiniNDN [22] and PX4 SITL [21] can exercise network behavior and application/autopilot integration respectively; neither establishes real-flight safety.

4.3 NDN-SF-DI: Inference Lifecycle, State, and Collaboration

4.3.1 Application Structure and Inference Workflows

NDN-SF-DI provides an inference application layer over the framework. It interprets model and task descriptions, coordinates model preparation and execution, and exposes inference results to callers. Its native requester and Provider interfaces include request submission, service

active and record locally even when live streaming is off. The recorded object is represented as a large-data reference to signed and encrypted Data segments. The ground station discovers available recordings through an authorized service invocation, then fetches the referenced segments by name and decrypts locally. Thus storage does not need to understand service authorization; authorization and key material are handled by NDNSF service messages, while replay or stale-data problems are handled by session identifiers, stream identifiers, and signed Data names.

4.6 Flight Control and Safety

Flight-control commands are naturally targeted at a specific drone. The proposed work uses Targeted NDNSF invocation for commands such as arm, takeoff, land, goto, emergency stop, and manual control. The drone-side adapter supports mock mode for MiniNDN tests and MAVLink/PX4 mode for SITL or real flight controllers [22, 23]. Manual-control commands are continuous control messages rather than mission services; their responses can carry telemetry so the GS can update altitude, battery, GPS, speed, controller state, and flight-safety indicators while the operator is flying.

Safety mechanisms include:

- readiness gates based on heartbeat, GPS/EKF, battery, mode, and arming checks;
- command acknowledgments from the flight controller;
- manual-control timeouts that stop sending stale control values;
- emergency stop or hold/land behavior after communication failure;
- explicit warnings that real propeller tests require a validated safety procedure.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

4.7 Multi-UAV Patrol

The proposed patrol service demonstrates provider collaboration. A ground station creates a waypoint mission and partitions it among two or more drones. The design clusters waypoints according to the number and position of available drones so each drone receives a compact region and returns to its own departure point after completing its segment. Each drone receives a mission segment, executes it through MAVLink, and returns status. If a drone fails or returns an incomplete result, the GS can issue a compensation request for the missing segment rather than restarting the entire mission.

This design illustrates a key NDN_SF principle: compensation requests are application-level parts of the same logical mission, while NDN_SF provides the secure service invocation, provider selection, and message verification substrate.

The current UAV-APP is a research prototype, not yet a field-ready or safety-certified ground-control system. This distinction is important for the proposal: practical use, SITL, and selected real-machine testing may expose new requirements or changes in command handling, telemetry freshness, video behavior, safety policy, deployment, or operator workflow. Those discoveries are part of the research value because they can feed back into NDN_SF’s framework mechanisms.

4.8 Evaluation Plan

The UAV application will be evaluated in three stages:

- 1. **MiniNDN regression tests:** controller, one or more drones, operator console behavior, video start/stop, encrypted recording playback, and patrol service execution.
- 2. **SITL tests:** PX4/jMAVSim control for arm, takeoff, land, manual control, and mission execution.
- 3. **Real-machine validation:** one PC and one or more drones with NFD routes, certificates, camera devices, and MAVLink connections.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Metrics will include command latency, video startup time, decoded frame rate, recording retrieval time, mission completion rate, compensation overhead, telemetry freshness, and failure recovery time.

This chapter supports RQ4 by defining UAV as the command-sensitive, real-time, mobile workload used to validate NDN SF beyond core framework tests. It also feeds back into RQ2 and RQ3 by using UAV behavior to stress Targeted invocation, provider collaboration, timeout diagnosis, and large-data references.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

Chapter 5

Distributed AI Inference over NDN

5.1 Motivation

Distributed AI inference is another strong test case for NDN. AI models are large, runtime artifacts are sensitive, intermediate tensors can be large, and inference providers may be heterogeneous. A model may be split across edge devices, drones, ground stations, GPUs, and cloud resources. The user wants to invoke an inference service by model or task name, while the system determines where model shards and stages should run.

For the dissertation, distributed inference stresses provider collaboration in a different way from UAV workloads. Providers are not collaborating because they are vehicles in a shared physical mission; they collaborate because a single inference task may require multiple model chunks, runtime artifacts, activation exchanges, and dependency-ordered execution steps. This workload therefore tests whether NDN can keep collaboration generic: the framework should coordinate authorized providers, large-data references, status diagnosis, and result verification, while the inference layer supplies the model-specific dependency graph and execution plan.

In IP-based systems, distributed inference often relies on custom orchestration: host addresses for model stages, separate object stores for model artifacts, custom integrity checks, ad hoc retry behavior, and manual placement policies. This makes it hard to build applications that remain correct when providers join, leave, fail, or move.

5.2 Approaches to Distributed AI Inference

The dissertation will review the main distributed inference approaches:

- **Replication/data parallel inference:** multiple providers hold the same model and serve independent requests.
- **Pipeline partitioning:** model stages execute sequentially on different providers,

when retrieval recovery behaves correctly.

4.2.3 Representative Collaboration: Multi-View Collection and Analysis

An operator chooses a target location or area on a map. Request Data carries only the service-protected descriptor needed to assess coverage, viewing feasibility, sensors, and timing. Eligible UAVs return positive ACKs. The proposed workflow assigns complementary viewing roles and a processing role, then delivers execution details and appropriate keys through Selection.

Each sensing UAV signs encrypted ViewImage Data under its own authorized namespace, with observation time and available pose/calibration metadata. The compute Provider retrieves permitted images, performs the specified analysis, and publishes signed Annotation Data. Its encrypted result supplies original image references and permitted image keys to the requester, which retrieves the images and renders annotations. The workflow reuses the original signed images rather than republishing them under another producer’s identity.

This case examines role assignments, named dependencies, result reuse, and incomplete-view or processing failures. It does not replace the control, telemetry, video, and recording validation above. GPS alone does not establish camera visibility, depth, or correspondence between views. Improved recognition or localization requires ground truth, a single-view control, and a specified fusion method. MiniNDN [22] and PX4 SITL [21] can exercise network behavior and application/autopilot integration respectively; neither establishes real-flight safety.

4.3 NDN-DI: Inference Lifecycle, State, and Collaboration

4.3.1 Application Structure and Inference Workflows

NDN-DI provides an inference application layer over the framework. It interprets model and task descriptions, coordinates model preparation and execution, and exposes inference results to callers. Its native requester and Provider interfaces include request submission, service

following split-computing ideas such as DNN partitioning across mobile, edge, and cloud resources [24, 25].

- **Layer or operator partitioning:** the model graph is split at finer granularity according to compute and memory constraints.
- **Early-exit inference:** computation can stop early when a confidence threshold is reached [26].
- **Collaborative inference:** multiple providers cross-check, aggregate, or vote on results.
- **Edge-cloud split inference:** lightweight preprocessing happens at the edge while heavier stages run on ground-station or cloud nodes.

This proposal uses two split terms to avoid ambiguity. A *stage split* divides the model into sequential stages, such as backbone, head, and merge tasks. This is the natural first choice for large language models and many pipeline-style models because each provider can own a contiguous part of the execution. A *shard split* divides work inside one stage across multiple providers. Sharding can help when no single provider can fit or execute a stage, but it usually adds communication and merge overhead. The planner should therefore prefer reusable stage layouts and introduce shard layouts only when provider memory or compute capacity requires them.

5.3 NDNFS-Based Inference Architecture

The proposed NDNFS-DistributedInference application treats an inference task as a named service with associated roles, stages, artifacts, tensor dependencies, and runtime requirements. A service policy defines:

- the service name, such as `/AI/YOLO/SplitInference`;
- model stages, chunks, and roles;

registration, result/event handling, and cancellation; conversation-related interfaces also carry continuation state. The presence of an interface is not evidence that every model or recovery configuration has passed end-to-end tests.

A single-Provider inference request exercises service discovery, permission checks, selection, protected input delivery, and result retrieval without requiring model partitioning. Preparing a cold model and reusing a resident model exercise different readiness and artifact-access paths. Model files and recorded artifacts may use managed repository storage; the forwarding Content Store remains opportunistic caching, not guaranteed persistence. DI defines model identity, compatibility, and placement policy, while NDNFS supplies service lifecycle and named-object exchange.

Streamed inference and conversation workflows extend coverage to incremental outputs, cancellation, terminal status, and request-specific state. Model weights, activation tensors, and key/value (K/V) state have distinct lifetimes. The model computes K/V tensors and the inference runtime retains them for later decoding; shared buffers can avoid copying the whole cache at each step [23]. Resident weights do not replace this state. DI owns tokenization and KV compatibility rules; framework validation concerns the authorized request, data exchange, and lifecycle through which the application uses that state. A KV speedup is not automatically an NDNFS speedup.

4.3.2 Representative Collaboration: Multi-Provider Inference

A planner combines a supported model description with runtime ACK capabilities, assigns Providers, and prepares their model fragments. Pipeline partitioning assigns sequential layers or layer groups. Tensor partitioning divides operations or tensors and requires the corresponding collective exchanges. Providers exchange activations or collective inputs over declared dependencies. ONNX [24] describes the computation graph; export alone does not implement cross-Provider collectives or establish numerical correctness. Model partitioning has substantial prior work [25, 26]; the NDNFS question concerns runtime assignment and authorized named-data execution.

- required artifacts such as ONNX chunks, PyTorch modules, or runtime artifacts;
- input and output contracts that describe tensor names, shapes, dtypes, and codecs;
- dependency edges that state which role produces each intermediate tensor bundle and which roles consume it;
- candidate providers, provider roles, and fallback policy.

The user should not need to manually create a plan for every request. Instead, the service controller or model-splitting tool should generate the plan when the model is prepared. For ONNX models, which provide a standard computation graph representation for machine-learning models [27], the proposed path is shown in Figure 5.




Figure 5: Automatic generation path from an ONNX computation graph to an NDNF-DistributedInference collaboration plan.

The same policy interface should also support non-ONNX models by accepting a model-specific splitter output. At runtime, the user invokes the service name; providers assigned to roles fetch the artifacts associated with their roles, prefetch predictable input tensor bundles,

Autoregressive generation also uses each generated token in a later decoding step. This temporal dependency differs from forward computation through the layers. The evaluation separates model preparation, prefill, time to first token, subsequent decoding, state storage, and data transfer. Claims remain specific to tested models, precisions, partitions, and devices. A feasible single-Provider baseline is retained; if a model cannot fit one device, distributed feasibility is meaningful but a speedup against a non-running baseline is not.

4.3.3 From a Model Graph to Runtime Assignments

Model preparation and Provider assignment occur at different times. Preparation identifies a supported model, its tensor interfaces, and candidate partitions. It can produce a reusable template describing fragments and their dependencies before any particular invocation. Runtime planning then uses the request and collected positive ACKs to bind those roles to actual Providers. A manually authored template is compatible with this design: automatic partition search is not required to test the protocol. The Controller’s authorization role does not make it the model partitioner.




Figure 3: Proposed separation of model preparation and request-specific collaboration. Arrows show processing order, not a claim that every model graph or assignment strategy is implemented.

execute their model chunk, publish output tensor bundles, and return or forward the final prediction.

5.4 Automatic Collaboration Plan Generation

A major proposed contribution is automatic generation of collaboration plans from AI model structure and resource information. The input includes:

- model graph, tensor dependency DAG, candidate cut points, and estimated activation sizes;
- provider compute, memory, accelerator, and battery state, with a homogeneous-provider default when detailed profiles are unavailable;
- network latency, RTT, and throughput estimates;
- artifact placement and cache state;
- security and permission constraints.

The output is an NDNF service plan that assigns model roles and chunks to providers, identifies required artifacts and tensor bundles, records predictable object names for prefetch, and defines fallback behavior.

5.5 Dependency-Driven Execution

The planned runtime should hide dependency management from AI application developers. A provider should not contain YOLO-specific or fully-connected network-specific communication code. Instead, each provider role reads its local input edges from the policy, prefetches all predictable input objects, merges tensor bundles, executes its assigned ONNX or model-specific chunk, and publishes the tensor bundles required by its output edges. This executor model is intended to support chain, fan-in, and fan-out graphs. YOLO 2x2 and small fully connected examples are therefore validation cases for the same dependency-driven execution design rather than separate hand-written systems.

Each fragment contract needs defined input and output tensor names, shapes or supported dynamic-shape constraints, element types, serialization, and model-version compatibility. A split across a skip connection or a branching graph may have multiple live outputs, not a single vector. The receiving fragment must obtain all required inputs before execution. A useful negative test supplies a correctly signed tensor with the wrong shape or model identity: protocol provenance may be valid while the application contract is not.

Pipeline partitioning assigns consecutive computation stages and transfers their boundary tensors. Tensor parallelism splits an operation across participants and adds exchanges such as reductions or gathers. Model replication assigns independent work to complete copies. These are different execution strategies, not interchangeable names for placing a model on several devices. Early exit changes how much model computation is performed and must be evaluated separately from placement. ONNX export supplies a graph representation [24]; usable cuts and collective operations still depend on the supported runtime and model operators.

Preparation cost matters when a Provider lacks the required artifact or compatible loaded session. Evaluation should distinguish artifact transfer, session creation, accelerator memory allocation, and execution. A warm run that reuses resident weights cannot be compared to a cold baseline without reporting that difference. Placement can favor a Provider that already holds compatible state, but that policy must respect request authorization and cannot assume that cached state is transferable to any new request.

4.3.4 Tensor Exchange and Autoregressive State

For a pipeline, an upstream fragment produces an activation tensor, makes it available as protected named Data, and supplies its reference through the declared dependency. The downstream fragment retrieves and validates that object, checks the tensor contract, and performs its assigned computation. A fan-out reuses an upstream output for several authorized consumers. A fan-in waits for the set of inputs required by the computation. The collaboration protocol decides which producers and edges are permitted; DI decides whether the tensor values and

5.6 Current Prototype Status

The current prototype has moved from deterministic trace execution toward real ONNX execution over the full MiniNDN network harness. The latest NativeTracer campaign uses a small Qwen-style model artifact and runs the same Backbone/Head/Merge workflow through real model execution, repo-backed artifact fetching, service manifests, and NDNsf large-data references. Providers execute assigned roles from the generated plan and exchange named intermediate tensor objects instead of embedding large tensors in service payloads.

The current layout campaign compares two executable runtime layouts under the same full-network harness. In ten successful runs per layout, the shared-backbone layout had mean latency 259.4 ms, p50 267.7 ms, p95 283.9 ms, and standard deviation 22.1 ms. The single-provider layout had mean latency 186.8 ms, p50 191.7 ms, p95 206.3 ms, and standard deviation 20.0 ms. For this small model, single-provider execution is faster because it avoids the fixed cost of cross-provider dependency exchange. This is a useful negative result: NDNsf-DI should not split a model merely because multiple providers are available; the planner must compare compute cost, intermediate-data size, provider capacity, RTT, bandwidth, and artifact placement before selecting a layout.

Earlier YOLO-style and small fully connected examples remain useful validation cases for dependency-driven execution, but the remaining research work is now to make the planner and runtime more general: remove model-specific assumptions from provider logic, use ONNX tensor DAGs and provider/network profiles to generate reusable plans, and evaluate when multi-provider execution helps rather than only adding coordination overhead.

5.7 Prototype and Evaluation Plan

The prototype will focus on ONNX split inference with both YOLO-style and small fully connected or transformer-style models, then generalize from chain-like examples to fan-in and fan-out dependency graphs. Planned experiments include:

layouts are suitable for the model.

Autoregressive inference adds a dependency across decoding steps. When K/V caching is enabled, computation after prefill uses the preceding token and retained K/V state instead of recomputing the full prefix. With layer-based partitioning, each participating runtime can retain the K/V tensors for its assigned layers. A generated token or the information needed to select it returns to the component driving the next decoding step. Consequently, the forward path within one evaluation can follow a directed acyclic graph while the whole generation process repeatedly traverses that graph over time.

Keeping weights in accelerator memory does not eliminate the logical K/V inputs. A runtime may reuse device buffers and bind them directly rather than serialize and transmit the entire cache on each step [23]. Reassignment or conversation continuation needs an explicit compatibility and state-transfer policy; silently starting a replacement with unrelated K/V state is not valid recovery. The evaluation separates savings from local K/V reuse, remote transfer costs, and framework orchestration costs.

The initial numerical oracle should compare a supported partitioned execution with an unsplit execution of the same model, precision, inputs, and decoding policy. Appropriate tensor tolerances or deterministic output comparisons must be specified before measurement. A small synthetic model can validate a tensor route, but cannot establish memory feasibility, throughput, or recovery for a large pretrained model. These controls make it possible to distinguish a working communication path from a correct and useful distributed-inference application.

4.4 Coverage, Evidence, and Completion Boundaries

Table 6 groups the applications by the framework mechanisms they exercise. It is a coverage plan, not a claim that every listed path is qualified.

Each evaluated row must identify the concrete application function, production NDNsf API path, configuration, test oracle, and durable evidence. Source implementation, focused tests, multi-process tests, and measured application outcomes are separate evidence levels. Local calls

- local baseline inference versus NDNST split inference;
- manually written policy versus automatically generated policy;
- single-provider layouts versus shared-backbone and multi-provider layouts;
- chain execution versus fan-in/fan-out dependency execution;
- cold-start artifact fetching versus cached artifacts;
- one-provider, two-provider, and multi-provider execution;
- planner decisions under different provider memory, compute, RTT, and bandwidth profiles;
- provider failure and compensation;
- consistency of distributed output relative to local inference;
- impact of artifact placement on startup latency.

Metrics will include end-to-end latency, artifact retrieval time, provider utilization, success rate, output consistency, recovery time, p50/p95 latency, latency variance, and planner prediction error.

This chapter supports RQ5 by defining distributed inference as the artifact-heavy, dependency-driven workload used to validate NDNST under multi-stage compute/data collaboration. It also feeds back into RQ2 by testing whether large-data references, provider roles, generated plans, and long-running selection status can remain generic rather than model-specific.

Table 6: Application workflows and intended NDNST validation coverage.

Application workflow	Framework mechanism	Required observation
UAV control/telemetry; DI single-Provider inference	Discovery, selection, service authorization	Correct Provider/request association; unauthorized invocation rejected
UAV live video; DI incremental output	Stream or invocation-output lifecycle	Validated outputs, termination, cancellation, and defined gap handling
UAV recordings; DI model/input objects	Named-object references and retrieval	Exact permitted object reconstructed; tampered or wrong object rejected
UAV mission status; DI preparation/conversation	Operation and request lifecycle	Correct status ordering; late actions cannot revive closed work
UAV multi-view; DI partitioned inference	Role assignment and named dependencies	Assigned producers only; missing dependencies and replacement handled
Permission changes in either application	Controller status and invocation checks	Grant discovered/applied; withdrawn permission rejected after update

cannot substitute for network tests, and a successful service exchange cannot establish correct image analysis, numerical output, or safe flight.

The current native implementation has focused and selected cross-process evidence, but the claimed model, device, and recovery configurations are not all qualified. Historical small-model measurements do not validate newer tensor/hybrid, multi-accelerator, or large-language-model recovery paths. The final campaign should cover the claimed mechanism categories using representative application workflows; it need not turn every optional interface or application feature into a separate dissertation completion gate. The broader application scope is preserved while each research conclusion remains tied to the evidence actually obtained.

- **user authorization:** whether the user may request a service;
- **provider authorization:** whether a provider may serve a service;
- **message integrity:** whether a service message is signed and valid;
- **data confidentiality:** whether the payload is encrypted for authorized participants;
- **replay protection:** whether one-time tokens have already been consumed.

3.8 Completed Foundation: NDN SF v0.1 Evaluation

The completed NDN SF evaluation compares NDN SF against NSC and gRPC in MiniNDN. The evaluation focuses on service latency and success rate under wired, lossy, intermittent-provider, wireless mobility, adaptive-admission, selective ACK, and provider-selection scenarios. The key result is that NDN SF can outperform NSC and improve reliability under provider diversity and intermittent availability, while maintaining reasonable latency relative to IP-based RPC in lossless conditions.

The strongest evidence comes from scenarios where endpoint binding or a single provider becomes fragile. In the lossy wired Memphis–CSU–UCLA path, packet round-trip loss is $1 - (1 - \ell)^4$ for per-link loss ℓ . NDN SF completed 599 of 600 requests at 1% link loss, where the packet round-trip loss is 3.94%, and 598 of 600 requests at 3% link loss, where the packet round-trip loss is 11.47%. At 10% link loss, where packet round-trip loss rises to 34.39%, the best measured NDN SF result was 539 of 600 requests. This identifies the bounded-time delivery limit of SVS-based service invocation under heavy loss, rather than a failure of the NDN SF service model. In a wired multi-provider experiment with intermittent provider rejection, NDN SF maintained 99.46% success at 20% provider-failure probability and 96.70% success at 30%, whereas the single-provider baselines dropped to roughly two-thirds success or lower. In the wireless mobility experiment, NDN SF achieved 96.25% success at 100 m AP range and 100% at larger ranges, while NSC and gRPC were much lower when bound to one mobile provider.

Chapter 5

Preliminary Evidence and Proposed Evaluation

5.1 Measured Findings Are Configuration-Specific

Invocation campaigns below measure discovery, selection, and execution under their recorded configurations; they do not test a multi-role collaboration plan. The earlier DI diagnostic is separate application evidence, not proof that the initial framework supported role-aware collaboration.

The following numbers are retained from recorded campaigns, not new experiments in this revision [27]. The independent unit for comparative uncertainty is the seed/run, not each correlated request. Success-conditioned latency must be read with completion and timeout rates. A baseline label identifies the evaluated harness configuration, not every feature of the framework named in that label.

Table 7: Selected evidence and limits; values are unchanged from their records.

Campaign	Observation	Supported interpretation / limit
Ten-seed switching holdout	1,312 matched switching requests; NDN SF minus gRPC seed-p95 difference: −597 ms, 95% CI [−992, −200]; versus NSC: −2618 ms, CI [−2916, −2315]	Lower conditional tail latency in the tested four-Provider configuration, not universal mobility or success-rate superiority
Six-seed work-efficiency repeat	NDN SF: 1,798/1,800 completed, 1,800 execution starts; gRPC-PAR-4: 1,800/1,800, 7,180 starts	Fewer starts than this parallel policy; starts are not messages, bytes, CPU time, or energy
Earlier small-model DI diagnostic	Ten successful runs/layout: mean 259.4 ms distributed versus 186.8 ms single Provider; p95 283.9 versus 206.3 ms	Splitting was slower for this model/configuration; small success-conditioned sample does not establish a general tail-latency result

In the work-efficiency comparison, NSC-4 also completed 1,800/1,800 requests, with 1.002 execution starts per request (rounded). This does not establish a work-efficiency advantage over NSC. The mobility comparison uses common traces, publication times, deadlines,

The adaptive-admission and Selective ACK experiments further show why the framework needs runtime metadata: admission control kept P95 latency below 270 ms under overload, and Selective ACK raised success from 1.82% to 100% when the nearest provider was overloaded. Together, these results support the v0.1 claim that one-to-many invocation, provider metadata, and explicit selection improve reliability without giving up data-centric security.

Table 5: Selected reliability results summarized from the submitted NDNNSF v0.1 evaluation.

Scenario	Condition	NSC	gRPC	NDN NSF
Intermittent wired providers	10% failure probability	83.25%	83.35%	99.82%
Intermittent wired providers	20% failure probability	66.58%	66.69%	99.46%
Intermittent wired providers	30% failure probability	49.92%	65.36%	96.70%
Wireless mobile providers	100 m AP range	45.42%	45.23%	96.25%
Wireless mobile providers	150 m AP range	68.75%	68.05%	100.00%
Wireless mobile providers	200 m AP range	88.33%	87.14%	100.00%

and four Providers per system. The switching holdout uses seeds 72–81 and a registered reachability-based classification. Of its 1,312 aligned requests per system, NDNNSF completed 1,311, gRPC 1,297, and NSC 1,296. Each seed-p95 is calculated from that system’s successful requests, so the latency comparison is affected by different failure sets. Overall coverage controls using seeds 62–71 are a separate campaign, not the parent sample of this holdout. The small-model diagnostic is retained from earlier records; its raw-run/model provenance remains to be rechecked. Communication/setup overhead is a possible explanation, not an established cause.

The transition experiment reports 357/357 completions for NDNNSF without pre-registering Provider addresses, 0/357 for static-list baselines, and 357/357 for pre-registered baseline variants. It isolates the discovery configuration. A gRPC resolver or other discovery integration can update a candidate list [28]; the result does not prove that gRPC or NSC cannot discover a new Provider.

Historical v0.1 loss results include 599/600 completions at 1% per-link loss, 598/600 at 3%, and a best recorded 539/600 at 10%. The last is a selected historical run, not a current aggregate estimate or a matched comparison against later baseline data. The expression $1 - (1 - \ell)^4$ assumes four independent packet traversals with equal loss ℓ ; it is not the full service-failure probability after retries. Timeouts under heavy loss do not, without further interventions, identify one exclusive cause or rule out protocol/policy effects.

Table 6: Selected mechanism results from the submitted NDNSF v0.1 evaluation.

Mechanism	Baseline Behavior	NDNSF Behavior
Adaptive admission under 100 RPS offered load	Without admission control, P95 latency reached 4254.39 ms and timeout rate reached 3.55%.	With admission control, admitted rate was limited to about 81 RPS, P95 latency stayed at 268.12 ms, and timeout rate was 0.00%.
Selective ACK with nearest provider overloaded	ACK-all selected the nearest overloaded provider for 2209 of 2250 requests and achieved 1.82% success.	Selective ACK suppressed overloaded ACKs, shifted selection to available providers, and achieved 100.00% success.
Real-machine intermittent availability	Single-provider NSC/gRPC stayed near 50% success when the chosen provider was unavailable for half the run.	NDNSF selected among two providers and reached 74.62% success at 10 RPS and 73.73% at 30 RPS, close to the expected availability window.

3.9 Limitations of NDNSF v0.1

NDNSF v0.1 establishes feasibility, but it does not yet fully address the requirements of dynamic service applications beyond controlled framework tests. The main limitations that motivate the remaining dissertation work are:

- **Low-latency known-provider control:** the normal Request/ACK/Selection/Response workflow is appropriate for provider discovery, but control commands to a known drone or device need a faster Targeted path while preserving authorization and replay protection.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

4.7 Multi-UAV Patrol

The proposed patrol service demonstrates provider collaboration. A ground station creates a waypoint mission and partitions it among two or more drones. The design clusters waypoints according to the number and position of available drones so each drone receives a compact region and returns to its own departure point after completing its segment. Each drone receives a mission segment, executes it through MAVLink, and returns status. If a drone fails or returns an incomplete result, the GS can issue a compensation request for the missing segment rather than restarting the entire mission.

This design illustrates a key NDN SF principle: compensation requests are application-level parts of the same logical mission, while NDN SF provides the secure service invocation, provider selection, and message verification substrate.

The current UAV-APP is a research prototype, not yet a field-ready or safety-certified ground-control system. This distinction is important for the proposal: practical use, SITL, and selected real-machine testing may expose new requirements or changes in command handling, telemetry freshness, video behavior, safety policy, deployment, or operator workflow. Those discoveries are part of the research value because they can feed back into NDN SF’s framework mechanisms.

4.8 Evaluation Plan

The UAV application will be evaluated in three stages:

- 1. **MiniNDN regression tests:** controller, one or more drones, operator console behavior, video start/stop, encrypted recording playback, and patrol service execution.
- 2. **SITL tests:** PX4/jMAVSim control for arm, takeoff, land, manual control, and mission execution.
- 3. **Real-machine validation:** one PC and one or more drones with NFD routes, certificates, camera devices, and MAVLink connections.

5.2 Evidence Needed for Each Research Question

Table 8: Claim-driven evaluation plan, not a completed-results checklist.

RQ	Comparison	Correctness / failure observation	Costs and outcomes
RQ1	ABE-backed authorization; DNMP-inspired role rules plus content encryption (Section 2.4.2); suitable group-key combinations	Unauthorized call never starts handler; replay, wrong challenge, grants, withdrawn/expired permissions, status delay	Completion, discovery delay, crypto work, credential/policy size, grant and withdrawal traffic
RQ2	Same task graph with static/dynamic assignment; controlled removal of each admission check	Wrong role, edge, name, key or execution; direct/cache parity; incomplete objects; cancellation or reassignment	End-to-end completion, validation cost, intermediate bytes, redundant work, bounded recovery
RQ3	Matched Provider counts, tasks, capacities, deadlines and failure schedules; feasible single-device control	Include failed and timed-out requests; distinguish network, admission, execution and result-retrieval failures	Completion and latency together; cold/warm costs; per-run uncertainty; conditions with no benefit

Each campaign must preserve its source/dependency/model identity, configurations, seeds, measurement start/end points, raw-run references, and analysis scripts. Scaling dimensions should follow the actual claim: candidates and services for discovery/authorization, edges and object sizes for collaboration. Ablation diagnoses a mechanism; an intentionally insecure variant is not a recommended deployment. A startup or collector failure must be diagnosed before attributing it to the harness or protocol. Preserve failed attempts; do not count an experiment that never reached its measurement boundary as a measured protocol failure rate.

5.3 Application Experiments and Interpretation

The experimental unit is a stated invocation or application session, with a recorded workload, software revision, and network configuration. Report submitted, admitted, selected, completed, rejected, timed-out, and cancelled work separately where the path exposes those outcomes. Success rate uses an explicit denominator and deadline. Latency distributions should

Metrics will include command latency, video startup time, decoded frame rate, recording retrieval time, mission completion rate, compensation overhead, telemetry freshness, and failure recovery time.

This chapter supports RQ4 by defining UAV as the command-sensitive, real-time, mobile workload used to validate NDNFSF beyond core framework tests. It also feeds back into RQ2 and RQ3 by using UAV behavior to stress Targeted invocation, provider collaboration, timeout diagnosis, and large-data references.

identify which requests they include, so faster completed requests do not hide a larger failure fraction. Repeated trials or independent seeds support variability estimates; a single exploratory run is reported as such.

For UAV sessions, record command completion latency, video start delay, useful frame delivery, recording retrieval completion, and status age at the operator. For multi-view work, additionally record the number of usable views, processing completion, and the analysis metric against labeled observations. The framework experiment first holds the recognition algorithm fixed. A separate application study is needed before attributing higher recognition accuracy to additional views.

For DI, report model preparation, prefill, time to first token, subsequent output rate, total completion, peak device memory, and bytes transferred between Providers. Compare cold and warm conditions separately. Evaluate a feasible single-Provider case alongside a supported partitioned case. If the distributed deployment exists because the model cannot fit one device, report that capacity benefit directly rather than inventing a speedup over an infeasible execution. Profile communication, serialization, cryptography, and model computation to locate the limiting cost.

Authorization evaluation should exercise both directions of service permission, not only successful decryption. Negative cases include an unpermitted caller, an unpermitted Provider, a validly signed message with the wrong protected value, and reuse after request state is closed. Grant experiments measure how an enrolled entity learns and applies a new permission. Withdrawal experiments measure update propagation, status validation, parameter/DKEY refresh, and rejection after the update is applied. Disconnected nodes and old ciphertext are separate cases with the limitations stated in Section 3.4.1.

The collaboration ablations retain request-ID isolation and signature validation. Removing assigned-producer checks tests a wrong-role object within the same request; removing edge checks tests an undeclared input from an otherwise authorized participant. Replacement tests inject old work after reassignment. This isolates the extension’s contribution instead

5.6 Current Prototype Status

The current prototype has moved from deterministic trace execution toward real ONNX execution over the full MiniNDN network harness. The latest NativeTracer campaign uses a small Qwen-style model artifact and runs the same Backbone/Head/Merge workflow through real model execution, repo-backed artifact fetching, service manifests, and NDNF large-data references. Providers execute assigned roles from the generated plan and exchange named intermediate tensor objects instead of embedding large tensors in service payloads.

The current layout campaign compares two executable runtime layouts under the same full-network harness. In ten successful runs per layout, the shared-backbone layout had mean latency 259.4 ms, p50 267.7 ms, p95 283.9 ms, and standard deviation 22.1 ms. The single-provider layout had mean latency 186.8 ms, p50 191.7 ms, p95 206.3 ms, and standard deviation 20.0 ms. For this small model, single-provider execution is faster because it avoids the fixed cost of cross-provider dependency exchange. This is a useful negative result: NDNF-DI should not split a model merely because multiple providers are available; the planner must compare compute cost, intermediate-data size, provider capacity, RTT, bandwidth, and artifact placement before selecting a layout.

Earlier YOLO-style and small fully connected examples remain useful validation cases for dependency-driven execution, but the remaining research work is now to make the planner and runtime more general: remove model-specific assumptions from provider logic, use ONNX tensor DAGs and provider/network profiles to generate reusable plans, and evaluate when multi-provider execution helps rather than only adding coordination overhead.

5.7 Prototype and Evaluation Plan

The prototype will focus on ONNX split inference with both YOLO-style and small fully connected or transformer-style models, then generalize from chain-like examples to fan-in and fan-out dependency graphs. Planned experiments include:

of claiming to solve cross-request confusion that the invocation foundation already handles. Correctness tests require the expected rejection and absence of prohibited handler execution, not merely an error message in a log.

Framework reuse can be assessed by documenting which invocation, authorization, object-transfer, and lifecycle components are shared by UAV and DI, and which adapters remain application-specific. Source size alone is not a sufficient maintainability metric. Record the changes needed to add a service type and the contracts that must be supplied. This supports the framework’s engineering rationale while leaving protocol correctness and performance to their own evidence.

Deployment records should include dependency versions, model hashes, topology, link parameters, clocks, retry policies, cache configuration, and key-management settings without publishing secrets. Separate application failures from startup, configuration, and collection failures. A failed setup is not evidence of a protocol’s inability to complete a request. These records allow the final dissertation to distinguish supported claims, useful negative results, and configurations that remain untested.

- local baseline inference versus NDNSE split inference;
- manually written policy versus automatically generated policy;
- single-provider layouts versus shared-backbone and multi-provider layouts;
- chain execution versus fan-in/fan-out dependency execution;
- cold-start artifact fetching versus cached artifacts;
- one-provider, two-provider, and multi-provider execution;
- planner decisions under different provider memory, compute, RTT, and bandwidth profiles;
- provider failure and compensation;
- consistency of distributed output relative to local inference;
- impact of artifact placement on startup latency.

Metrics will include end-to-end latency, artifact retrieval time, provider utilization, success rate, output consistency, recovery time, p50/p95 latency, latency variance, and planner prediction error.

This chapter supports RQ5 by defining distributed inference as the artifact-heavy, dependency-driven workload used to validate NDNSE under multi-stage compute/data collaboration. It also feeds back into RQ2 by testing whether large-data references, provider roles, generated plans, and long-running selection status can remain generic rather than model-specific.

No further page in this topic group.

Continue reading the other version,
or use the next topic bookmark.

4. **Evaluation and writing:** run final MiniNDN, SITL, selected real-machine experiments while converting the resulting evidence into dissertation chapters.

6.3 Evaluation Matrix

The final evaluation should connect the framework thesis to both application domains and show how the applications validate and extend NDN-SF mechanisms. Table 9 summarizes the planned evidence.

6.4 Timeline

6.5 Risks and Mitigation

- **Dependency risk:** UAV and AI applications depend on many libraries. Mitigation: controlled dependency environments and container-based tests.
- **Real flight risk:** physical UAV testing is safety-critical. Mitigation: start with MiniNDN and SITL, then use controlled real-machine tests without propellers before flight.
- **Performance risk:** security and sync overhead may affect latency. Mitigation: use Targeted invocation, local invocation, and large-data references as supporting mechanisms where the collaborative path would otherwise add avoidable overhead, together with sampled logging and careful benchmark profiles.
- **Scope risk:** UAV and AI inference are both large applications. Mitigation: implement minimal convincing prototypes first, then extend only the features needed to demonstrate the framework.

This chapter connects the research questions to the dissertation execution plan. It shows that each RQ has a corresponding evidence path, while leaving room for application-driven findings from UAV and DistributedInference to refine the final NDN-SF design.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Table 9: Evaluation matrix for the dissertation.

RQ	Dissertation Component	Evaluation Target	Evaluation Methodology
RQ1	Literature review	NDN service framework gap	Related-work comparison of RPC, robotics middleware, pub/sub, workflow systems, NDN service approaches, UAV networking, and distributed inference systems, summarized as framework requirements.
RQ2, RQ3	NDNSF core	Security, selection, targeted invocation, local composition, collaboration, and large-data support	MiniNDN regression and performance tests for permissions, NAC-ABE token routing, replay protection, provider selection, collaboration status, Targeted invocation, local invocation, and large-data references.
RQ4	UAV workload	Mobility, freshness, command lifecycle, media objects, and multi-UAV missions	GS/Drone MiniNDN tests, SITL MAVLink checks, command lifecycle logs, video stability, telemetry freshness, recording retrieval, mission progress, and failure recovery.
RQ5	DI workload	Model/data retrieval, split inference, dependency execution, planner decisions, and multi-provider collaboration	MiniNDN cluster tests for local-versus-distributed output equality, ONNX graph summary, generated policy checks, reference-based artifact provisioning, activation exchange, layout campaigns, provider utilization, and provider-failure cases.
RQ2, RQ4, RQ5	Operations	Reproducible deployment and selected real-machine validation	Certificate and route procedures, configuration checks, reproducible experiment setup, and selected real-machine tests.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Chapter 6

Research Plan and Timeline

6.1 Completed Work

The completed work includes the NDN service framework survey, NDN v0.1, core security design, provider selection, adaptive admission control, and preliminary MiniNDN evaluation. NDN v0.1 is not outside the dissertation; it is the completed first-stage core contribution on which the rest of the dissertation is built. Work after v0.1, including UAV-APP, DistributedInference, provider collaboration, and the supporting mechanisms needed to exercise it in realistic workloads, is treated as the application-driven extension and validation agenda to be completed and evaluated. Repository support may be used as an implementation component for large objects, but it is not presented as a separate dissertation thread.

6.2 Planned Work

- The remaining dissertation work is now primarily a stabilization, integration, and evaluation effort organized into four threads:
- Framework extensions:** finalize provider collaboration as the central NDN extension, and stabilize supporting mechanisms for low-latency Targeted invocation, trusted local composition, large-data references, and NAC-ABE operational support, with emphasis on correctness, documentation, and regression coverage.
 - UAV application:** stabilize the GS operator workflow, service-container organization, MAVLink/PX4 integration, live video, encrypted recording, object detection, and multi-UAV patrol as an application-driven validation of NDN.
 - Distributed inference:** complete the NativeTracer/ONNX execution path, generated service plans, reference-based artifact exchange, provider assignment, layout campaigns, and failure compensation, then freeze the evaluation scenarios.

Chapter 6

Research Plan, Feasibility, and Risks

6.1 Core Research Scope and Feasibility

The immediate priority is a frozen, runnable baseline for the final claims, not completion of every development feature. Select one UAV collection–processing graph and one supported DI partition first. Keep pipeline/tensor/hybrid generality, LLM streaming, and sophisticated recovery only where the implementation and evidence can support them. A negative performance result is an acceptable research finding when its conditions and costs are explained.

Table 9: Revised internal targets; defense is conditional on evidence and committee approval.

Period	Research milestone / decision
September 2026	Agree on three RQs and bounded claims; freeze supported application graphs and implementation/evidence baseline; resolve correctness blockers
October 2026	Execute registered correctness and matched comparison cases; reduce optional scope if an end-to-end gate remains open
November 2026	Freeze accepted results and provenance; write RQ answers, costs, and limitations; keep missing results visibly unverified
December 2026	Deliver an integrated dissertation draft to the advisor; prioritize analysis and revision over optional features
January 2027	Advisor/committee review, revisions, defense readiness decision, and confirmation of current university submission/notice requirements
February 2027	Target final dissertation defense after required reviews and evidence closure; allow time for substantive corrections
Spring 2027	Complete approved corrections and final submission by the applicable Graduate School deadlines; confirm exact dates separately

The February 2027 defense target depends on evidence and committee approval. CS guidance calls for an advisor-approved complete dissertation to reach committee members two weeks before the defense [29]. The Graduate School lists February 5, 2027 for the graduation application and doctoral candidacy form, and March 26, 2027 for the defended final ProQuest submission [30]; these dates must be rechecked before filing. If core end-to-end evidence is missing at the January review, scope or timing will be reconsidered with the committee.

4. **Evaluation and writing:** run final MiniNDN, SITL, selected real-machine experiments while converting the resulting evidence into dissertation chapters.

6.3 Evaluation Matrix

The final evaluation should connect the framework thesis to both application domains and show how the applications validate and extend NDN SF mechanisms. Table 9 summarizes the planned evidence.

6.4 Timeline

6.5 Risks and Mitigation

- **Dependency risk:** UAV and AI applications depend on many libraries. Mitigation: controlled dependency environments and container-based tests.
- **Real flight risk:** physical UAV testing is safety-critical. Mitigation: start with MiniNDN and SITL, then use controlled real-machine tests without propellers before flight.
- **Performance risk:** security and sync overhead may affect latency. Mitigation: use Targeted invocation, local invocation, and large-data references as supporting mechanisms where the collaborative path would otherwise add avoidable overhead, together with sampled logging and careful benchmark profiles.
- **Scope risk:** UAV and AI inference are both large applications. Mitigation: implement minimal convincing prototypes first, then extend only the features needed to demonstrate the framework.

This chapter connects the research questions to the dissertation execution plan. It shows that each RQ has a corresponding evidence path, while leaving room for application-driven findings from UAV and DistributedInference to refine the final NDN SF design.

6.2 Limits That Must Remain Visible

ABE adds authority, computation, and rekeying costs; disconnected revocation has propagation limits. Candidate ACKs do not prove future capacity or honest reporting. Cache availability does not guarantee task completion. Tests confined to one process do not establish distributed failure behavior. UAV safety and perception accuracy need evidence beyond service-message delivery. Findings from selected hardware and models must be interpreted within those configurations.

Table 9: Evaluation matrix for the dissertation.

RQ	Dissertation Component	Evaluation Target	Evaluation Methodology
RQ1	Literature review	NDN service framework gap	Related-work comparison of RPC, robotics middleware, pub/sub, workflow systems, NDN service approaches, UAV networking, and distributed inference systems, summarized as framework requirements.
RQ2, RQ3	NDNSF core	Security, selection, targeted invocation, local composition, collaboration, and large-data support	MiniNDN regression and performance tests for permissions, NAC-ABE token routing, replay protection, provider selection, collaboration status, Targeted invocation, local invocation, and large-data references.
RQ4	UAV workload	Mobility, freshness, command lifecycle, media objects, and multi-UAV missions	GS/Drone MiniNDN tests, SITL MAVLink checks, command lifecycle logs, video stability, telemetry freshness, recording retrieval, mission progress, and failure recovery.
RQ5	DI workload	Model/data retrieval, split inference, dependency execution, planner decisions, and multi-provider collaboration	MiniNDN cluster tests for local-versus-distributed output equality, ONNX graph summary, generated policy checks, reference-based artifact provisioning, activation exchange, layout campaigns, provider utilization, and provider-failure cases.
RQ2, RQ4, RQ5	Operations	Reproducible deployment and selected real-machine validation	Certificate and route procedures, configuration checks, reproducible experiment setup, and selected real-machine tests.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Table 10: Proposed dissertation timeline.

Period	Planned Work
June–July 2026	Stabilize NDNSE provider collaboration, its supporting Targeted/local/large-data mechanisms, and the UAV/DistributedInference prototypes.
August 2026	Complete UAV MiniNDN and SITL experiments for command lifecycle, telemetry readiness, live video, recording, map/mission workflows, and operational documentation.
September 2026	Complete DistributedInference integration tests, including ONNX graph summaries, generated policies, reference-based artifacts, dependency-driven execution, activation references, and provider-failure scenarios.
October 2026	Run final MiniNDN, SITL, and selected real-machine experiments; freeze experiment data, figures, tables, and experiment artifacts.
November 2026	Finish dissertation writing and revision, complete committee feedback cycles, and defend the dissertation in November 2026.

No further page in this topic group.

Continue reading the other version, or use the next topic bookmark.

Chapter 7
Conclusion

This proposal presents NDNFSF as a data-centric service framework for dynamic, secure, and collaborative networked applications. NDNFSF v0.1 demonstrates that NDN can support a generic service abstraction with authorization, provider selection, and adaptive behavior. The proposed work centers the next framework step on provider collaboration, supported by Targeted invocation, trusted local composition, and large-data references, and validates this direction through two demanding application domains: UAV networks and distributed AI inference.

The expected outcome is not merely another application implementation, but a framework and design methodology for building service-oriented systems where services, data, permissions, and collaboration plans are named and verifiable. UAV-APP and DistributedInference are important because they stress NDNFSF in different ways and reveal what the framework must provide, but the dissertation contribution remains the NDNFSF framework, its mechanisms, and the evidence that those mechanisms generalize beyond a single application.

References

[1] L. Zhang, A. Afanasyev, J. Burke, V. Jacobson, kc claffy, P. Crowley, C. Papadopoulos, L. Wang, B. Zhang, Named data networking, in: ACM SIGCOMM Computer Communication Review, 2014, pp. 66–73.

[2] C. Yi, A. Afanasyev, I. Moiseenko, L. Wang, B. Zhang, L. Zhang, A case for stateful forwarding plane, Computer Communications 36 (7) (2013) 779–791.

[3] T. Li, W. Shang, A. Afanasyev, L. Wang, L. Zhang, A brief introduction to ndn dataset synchronization (ndn sync), in: 2018 IEEE Military Communications Conference, 2018, pp. 612–618.

[4] P. Moll, V. Patil, L. Wang, L. Zhang, Sok: The evolution of distributed dataset synchronization solutions in ndn, in: Proceedings of the 9th ACM Conference on Information-Centric Networking, 2022, pp. 33–44.

[5] Statevectorsync: A library for distributed real-time applications over ndn, <https://github.com/named-data/ndn-svs>, , accessed 2026.

Chapter 7
Expected Outcomes and Proposal Summary

The research builds on NDNFSF’s initial ability to discover multiple candidates and select one Provider for an invocation. It evaluates that discovery/authorization foundation and investigates an extension for complementary Provider roles and named dependencies. The intended contribution is an evaluated service protocol and design rationale, not a new encryption primitive or the first use of named computation. RQ1 compares ABE-backed authorization with DNMP-inspired role rules paired with content encryption, including their management and withdrawal costs. RQ2 asks how runtime assignments become local rules for accepting named dependencies. RQ3 examines benefits and costs for invocation and collaboration separately. Existing campaigns establish limited invocation findings; the collaboration and cross-application claims still require their declared end-to-end evidence. The proposal therefore asks the committee to assess the research scope, comparison designs, validation criteria, and feasibility of the completion plan. The dissertation will report supported answers and limitations, including conditions in which the proposed mechanisms offer no advantage.

References

[1] L. Zhang, A. Afanasyev, J. Burke, V. Jacobson, kc claffy, P. Crowley, C. Papadopoulos, L. Wang, B. Zhang, Named data networking, in: ACM SIGCOMM Computer Communication Review, 2014, pp. 66–73.

[2] Named Data Networking Project, NDN Packet Format Specification: Data Packet, <https://docs.named-data.net/NDN-packet-spec/current/data.html>, , accessed September 9, 2026.

[3] Y. Yu, A. Afanasyev, D. Clark, k. claffy, V. Jacobson, L. Zhang, Schematizing trust in named data networking, in: Proceedings of the 2nd ACM Conference on Information-Centric Networking, 2015, pp. 177–186. doi:10.1145/2810156.2810170.

[4] T. Li, W. Shang, A. Afanasyev, L. Wang, L. Zhang, A brief introduction to ndn dataset synchronization (ndn sync), in: 2018 IEEE Military Communications Conference, 2018, pp. 612–618.

Chapter 7
Conclusion

This proposal presents NDNFSF as a data-centric service framework for dynamic, secure, and collaborative networked applications. NDNFSF v0.1 demonstrates that NDN can support a generic service abstraction with authorization, provider selection, and adaptive behavior. The proposed work centers the next framework step on provider collaboration, supported by Targeted invocation, trusted local composition, and large-data references, and validates this direction through two demanding application domains: UAV networks and distributed AI inference.

The expected outcome is not merely another application implementation, but a framework and design methodology for building service-oriented systems where services, data, permissions, and collaboration plans are named and verifiable. UAV-APP and DistributedInference are important because they stress NDNFSF in different ways and reveal what the framework must provide, but the dissertation contribution remains the NDNFSF framework, its mechanisms, and the evidence that those mechanisms generalize beyond a single application.

References

[1] L. Zhang, A. Afanasyev, J. Burke, V. Jacobson, kc claffy, P. Crowley, C. Papadopoulos, L. Wang, B. Zhang, Named data networking, in: ACM SIGCOMM Computer Communication Review, 2014, pp. 66–73.

[2] C. Yi, A. Afanasyev, I. Moiseenko, L. Wang, B. Zhang, L. Zhang, A case for stateful forwarding plane, Computer Communications 36 (7) (2013) 779–791.

[3] T. Li, W. Shang, A. Afanasyev, L. Wang, L. Zhang, A brief introduction to ndn dataset synchronization (ndn sync), in: 2018 IEEE Military Communications Conference, 2018, pp. 612–618.

[4] P. Moll, V. Patil, L. Wang, L. Zhang, Sok: The evolution of distributed dataset synchronization solutions in ndn, in: Proceedings of the 9th ACM Conference on Information-Centric Networking, 2022, pp. 33–44.

[5] Statevectorsync: A library for distributed real-time applications over ndn, <https://github.com/named-data/ndn-svs>, , accessed 2026.

Chapter 7
Expected Outcomes and Proposal Summary

The research builds on NDNFSF’s initial ability to discover multiple candidates and select one Provider for an invocation. It evaluates that discovery/authorization foundation and investigates an extension for complementary Provider roles and named dependencies. The intended contribution is an evaluated service protocol and design rationale, not a new encryption primitive or the first use of named computation. RQ1 compares ABE-backed authorization with DNMP-inspired role rules paired with content encryption, including their management and withdrawal costs. RQ2 asks how runtime assignments become local rules for accepting named dependencies. RQ3 examines benefits and costs for invocation and collaboration separately. Existing campaigns establish limited invocation findings; the collaboration and cross-application claims still require their declared end-to-end evidence. The proposal therefore asks the committee to assess the research scope, comparison designs, validation criteria, and feasibility of the completion plan. The dissertation will report supported answers and limitations, including conditions in which the proposed mechanisms offer no advantage.

References

[1] L. Zhang, A. Afanasyev, J. Burke, V. Jacobson, kc claffy, P. Crowley, C. Papadopoulos, L. Wang, B. Zhang, Named data networking, in: ACM SIGCOMM Computer Communication Review, 2014, pp. 66–73.

[2] Named Data Networking Project, NDN Packet Format Specification: Data Packet, <https://docs.named-data.net/NDN-packet-spec/current/data.html>, , accessed September 9, 2026.

[3] Y. Yu, A. Afanasyev, D. Clark, k. claffy, V. Jacobson, L. Zhang, Schematizing trust in named data networking, in: Proceedings of the 2nd ACM Conference on Information-Centric Networking, 2015, pp. 177–186. doi:10.1145/2810156.2810170.

[4] T. Li, W. Shang, A. Afanasyev, L. Wang, L. Zhang, A brief introduction to ndn dataset synchronization (ndn sync), in: 2018 IEEE Military Communications Conference, 2018, pp. 612–618.

[6] ndn-cxx: Ndn c++ library with experimental extensions, <https://github.com/named-data/ndn-cxx>, , accessed 2026.

[7] Nac: Name-based access control library for named data networking, <https://github.com/named-data/name-based-access-control>, , accessed 2026.

[8] grpc: A high performance, open source universal rpc framework, <https://grpc.io/>, , accessed 2026.

[9] R. T. Fielding, Architectural styles and the design of network-based software architectures, Ph.D. thesis, University of California, Irvine (2000).

[10] M. Quigley, K. Conley, B. Gerkey, J. Faust, T. Foote, J. Leibs, R. Wheeler, A. Y. Ng, Ros: an open-source robot operating system, ICRA Workshop on Open Source Software.

[11] Ros 2 documentation, <https://docs.ros.org/>, , accessed 2026.

[12] V. John, X. Liu, A survey of distributed message broker queues, arXiv preprint arXiv:1704.00411.

[13] C. Tschudin, M. Sifalakis, Named functions and cached computations, in: 2014 IEEE 11th Consumer Communications and Networking Conference, 2014, pp. 851–857.

[14] M. Krol, I. Psaras, Nfaas: Named function as a service, in: Proceedings of the 4th ACM Conference on Information-Centric Networking, 2017, pp. 134–144.

[15] M. A. ur Rehman, B.-S. Kim, Cfec: An ultra-low latency microservices-based in-network computing framework for information-centric iovs, IEEE Transactions on Services Computing 17 (6) (2024) 3199–3212.

[16] M. Krol, K. Habak, D. Oran, D. Kutscher, I. Psaras, Rice: Remote method invocation in icn, in: Proceedings of the 5th ACM Conference on Information-Centric Networking, 2018, pp. 1–11.

[17] T. Meirovitch, L. Zhang, Nsc–named service calls, or a remote procedure call for ndn, in: NDN Technical Report NDN-0074, Revision 1, 2021, pp. 1–20.

[18] K. Nichols, Lessons learned building a secure network measurement framework using basic ndn, in: Proceedings of the 6th ACM Conference on Information-Centric Networking, 2019, pp. 112–122.

[19] M. Imran, et al., Mia-ndn: Microservice-centric interest aggregation in named data networking, Sensors 23 (3) (2023) 1411.

[20] R. Tourani, et al., Towards security-as-a-service in multi-access edge, in: Proceedings of the 4th ACM/IEEE Symposium on Edge Computing, 2019, pp. 358–363.

[21] V. Goyal, O. Pandey, A. Sahai, B. Waters, Attribute-based encryption for fine-grained access control of encrypted data, in: Proceedings of the 13th ACM Conference on Computer and Communications Security, 2006, pp. 89–98.

[5] P. Moll, V. Patil, L. Wang, L. Zhang, Sok: The evolution of distributed dataset synchronization solutions in ndn, in: Proceedings of the 9th ACM Conference on Information-Centric Networking, 2022, pp. 33–44.

[6] C. Tschudin, M. Sifalakis, Named functions and cached computations, in: 2014 IEEE 11th Consumer Communications and Networking Conference, 2014, pp. 851–857.

[7] M. Krol, I. Psaras, Nfaas: Named function as a service, in: Proceedings of the 4th ACM Conference on Information-Centric Networking, 2017, pp. 134–144.

[8] M. Krol, K. Habak, D. Oran, D. Kutscher, I. Psaras, Rice: Remote method invocation in icn, in: Proceedings of the 5th ACM Conference on Information-Centric Networking, 2018, pp. 1–11.

[9] T. Meirovitch, L. Zhang, Nsc–named service calls, or a remote procedure call for ndn, in: NDN Technical Report NDN-0074, Revision 1, 2021, pp. 1–20.

[10] K. Nichols, Lessons learned building a secure network measurement framework using basic ndn, in: Proceedings of the 6th ACM Conference on Information-Centric Networking, 2019, pp. 112–122. doi:10.1145/3357150.3357397.

[11] grpc: A high performance, open source universal rpc framework, <https://grpc.io/>, , accessed 2026.

[12] M. Quigley, K. Conley, B. Gerkey, J. Faust, T. Foote, J. Leibs, R. Wheeler, A. Y. Ng, Ros: an open-source robot operating system, ICRA Workshop on Open Source Software.

[13] Ros 2 documentation, <https://docs.ros.org/>, , accessed 2026.

[14] V. John, X. Liu, A survey of distributed message broker queues, arXiv preprint arXiv:1704.00411.

[15] M. S. ud din, M. A. ur Rehman, B.-S. Kim, Cfec: An ultra-low latency microservices-based in-network computing framework for information-centric iovs, IEEE Transactions on Services Computing 17 (6) (2024) 3199–3212.

[16] M. Imran, et al., Mia-ndn: Microservice-centric interest aggregation in named data networking, Sensors 23 (3) (2023) 1411.

[17] R. Tourani, et al., Towards security-as-a-service in multi-access edge, in: Proceedings of the 4th ACM/IEEE Symposium on Edge Computing, 2019, pp. 358–363.

[18] C. Barrios, E. Alker, M. Kumar, CaSCON: Cache-assisted service composition orchestration leveraging NDN, in: 2025 IEEE International Conference on Smart Computing (SMARTCOMP), 2025, pp. 122–129. doi:10.1109/SMARTCOMP65954.2025.00061.

[22] A. Koubâa, A. Allouch, M. Alajlan, Y. Javed, A. Belghith, M. Khalgui, Micro air vehicle link (mavlink) in a nutshell: A survey, IEEE Access 7 (2019) 87658–87680. doi:10.1109/ACCESS.2019.2924410.

[23] Px4 autopilot user guide, <https://docs.px4.io/>, , accessed 2026.

[24] Y. Kang, J. Hauswald, C. Gao, A. Rovinski, T. Mudge, J. Mars, L. Tang, Neurosurgeon: Collaborative intelligence between the cloud and mobile edge, in: Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems, 2017, pp. 615–629. doi:10.1145/3037697.3037698.

[25] S. Teerapittayanon, B. McDanel, H. T. Kung, Distributed deep neural networks over the cloud, the edge and end devices, in: 2017 IEEE 37th International Conference on Distributed Computing Systems, 2017, pp. 328–339. doi:10.1109/ICDCS.2017.226.

[26] S. Teerapittayanon, B. McDanel, H. T. Kung, Branchynet: Fast inference via early exiting from deep neural networks, in: 2016 23rd International Conference on Pattern Recognition, 2016, pp. 2464–2469. doi:10.1109/ICPR.2016.7900006.

[27] Open neural network exchange (onnx), <https://github.com/onnx/onnx>, , accessed 2026.

[19] S. Dulal, T. Yu, S. Liu, A. R. Thieme, L. Zhang, L. Wang, Enhancing NAC-ABE to support access control for mHealth applications and beyond, arXiv:2311.07299 (2023). doi:10.48550/arXiv.2311.07299.
URL <https://arxiv.org/abs/2311.07299>

[20] A. Koubâa, A. Allouch, M. Alajlan, Y. Javed, A. Belghith, M. Khalgui, Micro air vehicle link (mavlink) in a nutshell: A survey, IEEE Access 7 (2019) 87658–87680. doi:10.1109/ACCESS.2019.2924410.

[21] Px4 autopilot user guide, <https://docs.px4.io/>, , accessed 2026.

[22] Minindn: A mininet-based ndn emulator, <https://github.com/named-data/mini-ndn>, , accessed 2026.

[23] ONNX Runtime Authors, How to configure the past present share buffer, <https://onnxruntime.ai/docs/genai/howto/past-present-share-buffer.html>, , accessed September 11, 2026.

[24] Open neural network exchange (onnx), <https://github.com/onnx/onnx>, , accessed 2026.

[25] Y. Kang, J. Hauswald, C. Gao, A. Rovinski, T. Mudge, J. Mars, L. Tang, Neurosurgeon: Collaborative intelligence between the cloud and mobile edge, in: Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems, 2017, pp. 615–629. doi:10.1145/3037697.3037698.

[26] S. Teerapittayanon, B. McDanel, H. T. Kung, Distributed deep neural networks over the cloud, the edge and end devices, in: 2017 IEEE 37th International Conference on Distributed Computing Systems, 2017, pp. 328–339. doi:10.1109/ICDCS.2017.226.

[27] NDN SF Project, Proposal research revision: Claim and evidence audit, Companion repository document: docs/PAPER/proposal-defense/research-revision-audit.md, , september 9, 2026; includes frozen campaign paths, evidence limits, and unresolved research gates.

[28] gRPC Authors, Custom name resolution, <https://grpc.io/docs/guides/custom-name-resolution/>, , accessed September 9, 2026.

[29] University of Memphis, Department of Computer Science, Information for computer science ph.d. students, https://www.memphis.edu/cs/pdfs/forms_grad_phd_programinfo.pdf, , accessed September 9, 2026.

[30] University of Memphis Graduate School, Graduate school calendar, <https://www.memphis.edu/gradschool/news-events/calendar.php>, , spring 2027 calendar, checked September 14, 2026.